

Estimation of obesity levels in individuals from Colombia, Peru and Mexico

Marwan Agourram (967349)

Accademic Year 2021-2022

Abstract The aim of this project is to analyse, develop predictions models and define clusters of the data "Obesity levels based on eating habits and physical condition in individuals from Colombia, Peru and Mexico" from UCI Machine Learning Repository¹.

The first part of the study is focused on developing prediction models using Supervised Learning techniques in order to predict the Obesity Level for each individual. The second part is focused on clustering and data reduction using Unsupervised Learning techniques.

Contents

1 Part 1: Supervised Learning	3
1.1 Introduction	3
1.1.1 Data Inspection	3
1.1.2 Data Cleaning and Pre-processing	5
1.2 Techniques and implemented algorithms	7
1.3 Linear Regression	8
1.3.1 Simple Linear Regression	8
1.3.2 Stepwise Regression	11
1.3.3 LASSO Regression	12
1.3.4 Linear Regression models comparison	15
1.4 Classification	17
1.5 Beyond Linearity	18
1.5.1 Decision Trees	18
1.5.2 Random Forest	22
1.6 Conclusions	24
2 Part 2: Unsupervised Learning	26
2.1 Multiple Correspondence Analysis	26
2.2 Hierarchical Clustering on Principal Components	30
2.3 Conclusion	34
A Code	36

¹Estimation of obesity levels based on eating habits and physical condition . (2019). UCI Machine Learning Repository.

List of Figures

1	Distribution of all numerical variables in the data.	4
2	Distribution of Age, Height and Weight.	5
3	<i>Left</i> : Boxplot with outliers included. <i>Right</i> : Boxplot without outliers.	6
4	Distribution of all the features, including the dummies, after one-hot and ordinal encoding was applied.	7
5	Training set and Test set dimensions.	8
6	Linear Regression model on all of the 19 features, with response variable NObeyesdad.	9
7	Linear Regression Model performance on different aspects.	10
8	Heteroscedasticity and Autocorrelation check on Linear Regression model.	10
9	Moderate correlations resulting from Autocorrelation check on Linear Regression model.	10
10	Stepwise Linear Regression model on a selected subset of features, with response variable NObeyesdad.	11
11	Stepwise Linear Regression model performance of different aspects.	12
12	Plot of the minimum $\log(\lambda)$ selected through cross validation on the training set	13
13	Selected features using LASSO approach.	13
14	LASSO Linear Regression model performance of different aspects.	14
15	Moderate correlations resulting from Autocorrelation check on LASSO Linear Regression model	14
16	LASSO Regression model on a selected subset of features, with response variable NObeyesdad.	15
17	Simple Linear Regression, LASSO Linear Regression and Stepwise Linear Regression comparison.	16
18	Simple Linear Regression, LASSO Linear Regression and Stepwise Linear Regression comparison using a Spider web plot.	16
19	Multinomial Logistic Regression call in Rstudio.	17
20	Complexity parameter for the Classification Tree.	20
21	Pruned Classification Tree with $cp = 0.020$	21
22	Random Forest Error with respect to the Number of Trees.	22
23	Random Forest include variables and relative importance.	23
24	Random Forest OOB and Test Error.	24
25	Frequency for the <i>Active Variables</i>	27
26	<i>Left</i> : Eigenvalues-Variance for MSE in each dimension. <i>Right</i> : Correlation between variables and principal dimensions.	28
27	Quality of representation of variable categories.	29
28	Quality of representation of individuals.	30
29	Dendrogram obtained after Hierarchical Clustering.	31
30	Factor map for the Hierarchical Clustering.	32

List of Tables

1 Part 1: Supervised Learning

1.1 Introduction

In this report I am going to use "*Dataset for estimation of obesity levels based on eating habits and physical condition in individuals from Colombia, Peru and Mexico*" paper from Mendoza F., et al. as a starting point. The paper presents data for the estimation of obesity level in individuals from Mexico, Perú and Colombia based on their eating habits and physical conditions.

1.1.1 Data Inspection

The data contains informations for the estimation of the obesity level in individuals from three different countries, with ages between 14 and 61 years old and different eating and physical conditions habits. The data was collected using a web platform with a survey where anonymous answered each question, specifying their age and gender and some physical attributes, such as height and weight.

These collected information end up consisting in a dataset with 17 attributes for 2,111 records. As mentioned above, four of these attributes are physical attributes for each individual:

- Gender: individual's gender (binary: "Male", "Female");
- Age: individual's age (numeric value: from 14 to 61);
- Height: individual's height in meters (numeric value: from 1.45 m to 1.98 m);
- Weight: individual's weight in kilograms (numeric value: from 39 kg to 173 kg).

The remaining 12 attributes are obtained from the survey's response to the reported questions:

- family_history_with_overweight: "Has a family member suffered or suffers from overweight?" (nominal: "yes", "no");
- FAVC: "Do you eat high caloric food frequently?" (nominal: "yes", "no");
- FCVC: "Do you usually eat vegetables in your meals?" (numeric: "1" - never, "2" - sometimes, "3" - always);
- NCP: "How many main meals do you have daily?" (numeric: "1" - between 1 or 2 meals, "2" - three meals, "3" - more than three meals);
- CAEC: "Do you eat any food between meals?" (nominal: "no", "sometimes", "frequently", "always");
- SMOKE: "Do you smoke?" (nominal: "yes", "no");
- CH2O: "How much water do you drink daily?" (numeric: "1" - less than 1 liter, "2" - 1 or 2 liters, "3" - more than 2 liters);
- SCC: "Do you monitor the calories you eat daily?" (nominal: "yes", "no");
- FAF: "How often do you have physical activity?" (numeric: "0" - no physical activity, "1" - 1 or 2 days, "2" - 2 or 3 days, "3" - 4 or 5 days);
- TUE: "How much time do you use technological devices such as cell phone, videogames, television, computer and others?" (numeric: "0" - 0-2 hours, "1" - 3-5 hours, "2" - more than 5 hours);
- CALC: "How often do you drink alcohol?" (nominal: "no", "sometimes", "frequently", "always");
- MTRANS: "Which transportation do you usually use?" (nominal: "automobile", "motorbike", "bike", "public transportation", "walking").

Then the class variable NObesdad was created using the Body Mass Index (BMI) equation as stated by the World Health Organization²

$$\text{BMI} = \frac{\text{Weight}}{\text{Height}^2}, \quad (1)$$

which values allows us to obtain the following nutritional status:

- Insufficient Weight: $\text{BMI} < 18.5$;
- Normal Weight: $18.5 \leq \text{BMI} \leq 24.9$;
- Overweight Level I and II: $25.0 \leq \text{BMI} \leq 29.9$;
- Obesity Type I: $30 \leq \text{BMI} \leq 34.9$;
- Obesity Type II: $35.0 \leq \text{BMI} \leq 39.9$;
- Obesity Type III: $\text{BMI} \geq 40$.

In Figure 1 we can observe the distribution of all the numeric variable, the number of missing values, and the range of the data with their respective mean and standard deviation.

```
> describe_distribution(data)
```

Variable	Mean	SD	IQR	Range	Skewness	Kurtosis	n	n_Missing
Age	24.32	6.36	6.00	[14.00, 61.00]	1.52	2.80	2111	0
Height	1.70	0.09	0.14	[1.45, 1.98]	-9.12e-03	-0.57	2111	0
Weight	86.59	26.19	42.06	[39.00, 173.00]	0.26	-0.70	2111	0
FCVC	2.42	0.58	1.00	[1.00, 3.00]	-0.43	-0.71	2111	0
NCP	2.62	0.73	0.00	[1.00, 3.00]	-1.56	0.67	2111	0
CH2O	2.01	0.69	0.00	[1.00, 3.00]	-0.02	-0.89	2111	0
FAF	1.01	0.90	2.00	[0.00, 3.00]	0.46	-0.71	2111	0
TUE	0.66	0.67	1.00	[0.00, 2.00]	0.52	-0.76	2111	0

Figure 1: Distribution of all numerical variables in the data.

Before going any further, I would like to highlight the fact that the physical features (Age, Height and Weight) have a large difference in variance and in the standard deviation compared to the other numerical variables. Here in Figure 2 we can have a look at these three variable distributions, represented with an histogram. In the following pages I will deal with these three variables.

²<https://www.who.int/europe/news-room/fact-sheets/item/a-healthy-lifestyle—who-recommendations>




Figure 2: Distribution of Age, Height and Weight.

1.1.2 Data Cleaning and Pre-processing

This report is based on a dataset which is composed by qualitative variables and quantitative ones. As we have seen while looking at the different features, most of the variables are responses based on multiple-choices answers, while some of the variables are a representation of the physical features of each individual who answered the survey.

In Figure 1 is evident that some of the numerical variables have different distributions with respect to all the other variables. In particular, in Figure 2 I displayed the distribution of variables that are more noticeable. The boxplot representations of these three variables are shown in Figure 3. It is easy to see that the variables Age and Weight have some outliers. In order to detect and remove these outliers, I applied the Interquartile Range rule. First, I computed the first and third quartile for Age and Weight, then I computed the lower limit (Equation 2) and upper limit (Equation 3) for both variables. An observation is considered to be an outlier when it exceeds the upper limit or it is lower than the lower limit.

$$LL = Q_1 - 1.5 * IQR, \quad (2)$$

$$UP = Q_3 + 1.5 * IQR \quad (3)$$

From computing the interquartile range, it results that 161 observations are considered to be outliers and, therefore, I proceeded to remove them from the dataset. Figure 3 shows the distribution for the three considered numerical variables after outliers have been removed. Therefore, I will mainly deal with observations coming from individuals whose age ranges between 14 and 35 years old, excluding the elder individuals. The rationale behind this choice is that the main goal of this part is to predict potential obesity problem in an individual based on their habits, and including these outliers may lead to some wrong conclusions. Also, I excluded individuals whose weight was exceeding the interquartile range. Once again, the rationale behind this decision is to ensure that the following conclusions are not affected by noise.

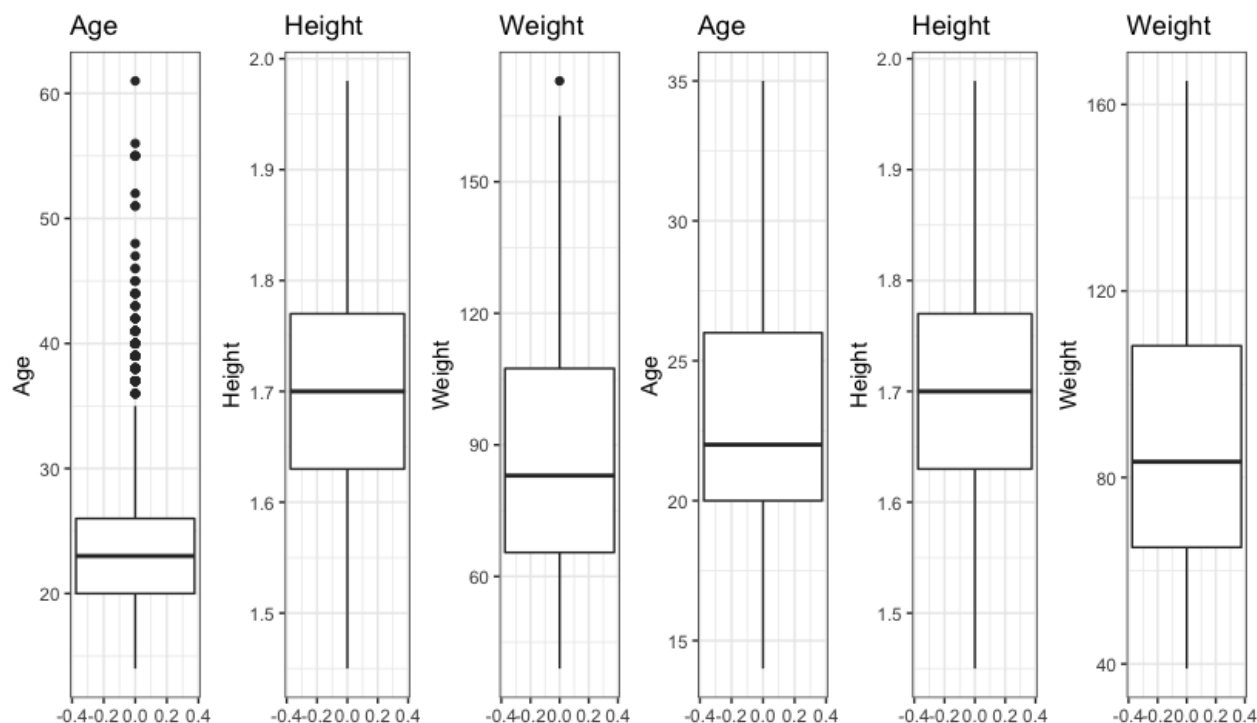


Figure 3: *Left*: Boxplot with outliers included. *Right*: Boxplot without outliers.

After removing the outliers observations, I proceeded to apply one-hot encoding technique on those categorical features that has a binary response class: Gender, family_history_with_overweight, FAVC, SMOKE, SCC, and MTRANS. One-hot encoding will generate dummy variables for each of the considered features. Then, I proceeded to apply ordinal encoding technique for the remaining categorical variables, that is, the ones which response classes were not binary: CAEC, CALC, and NObeyesdad.

Figure 4 reports the distribution of all the features in the dataset after applying previously mentioned techniques. It is worth to notice that now our data set consists of 26 features, due to one-hot encoding on binary categorical response class features, and 1950 observations, due to outliers detection and removal.

```
> describe_distribution(data_igr)
```

Variable	Mean	SD	IQR	Range	Skewness	Kurtosis	n	n_Missing
GenderFemale	0.49	0.50	1.00	[0.00, 1.00]	0.06	-2.00	1950	0
GenderMale	0.51	0.50	1.00	[0.00, 1.00]	-0.06	-2.00	1950	0
Age	23.00	4.38	6.00	[14.00, 35.00]	0.79	0.02	1950	0
Height	1.71	0.09	0.14	[1.45, 1.98]	6.95e-03	-0.52	1950	0
Weight	86.59	26.76	43.25	[39.00, 165.06]	0.23	-0.80	1950	0
family_history_with_overweightno	0.19	0.39	0.00	[0.00, 1.00]	1.58	0.48	1950	0
family_history_with_overweightyes	0.81	0.39	0.00	[0.00, 1.00]	-1.58	0.48	1950	0
FAVCno	0.12	0.32	0.00	[0.00, 1.00]	2.37	3.62	1950	0
FAVCyes	0.88	0.32	0.00	[0.00, 1.00]	-2.37	3.62	1950	0
FCVC	2.43	0.59	1.00	[1.00, 3.00]	-0.49	-0.66	1950	0
NCP	2.62	0.73	0.00	[1.00, 3.00]	-1.58	0.74	1950	0
SMOKEno	0.98	0.14	0.00	[0.00, 1.00]	-6.86	45.14	1950	0
SMOKEyes	0.02	0.14	0.00	[0.00, 1.00]	6.86	45.14	1950	0
CH2O	2.02	0.69	0.00	[1.00, 3.00]	-0.03	-0.89	1950	0
SCCno	0.95	0.21	0.00	[0.00, 1.00]	-4.22	15.84	1950	0
SCCyes	0.05	0.21	0.00	[0.00, 1.00]	4.22	15.84	1950	0
FAF	1.02	0.89	2.00	[0.00, 3.00]	0.44	-0.72	1950	0
TUE	0.70	0.68	1.00	[0.00, 2.00]	0.44	-0.81	1950	0
MTRANSAutomobile	0.16	0.37	0.00	[0.00, 1.00]	1.87	1.51	1950	0
MTRANSBike	3.08e-03	0.06	0.00	[0.00, 1.00]	17.96	320.83	1950	0
MTRANSMotorbike	4.62e-03	0.07	0.00	[0.00, 1.00]	14.63	212.22	1950	0
MTRANSPublic_Transportation	0.81	0.40	0.00	[0.00, 1.00]	-1.55	0.39	1950	0
MTRANSWalking	0.03	0.17	0.00	[0.00, 1.00]	5.70	30.56	1950	0

Figure 4: Distribution of all the features, including the dummies, after one-hot and ordinal encoding was applied.

Finally, some variables have been rescaled in order to let the models learn faster and to better perform in prediction, while others (such as Age, Height, Weight) were not scaled.

1.2 Techniques and implemented algorithms

In this first part I applied Supervised Learning techniques in order to analyse the data and generate prediction models in order to infer about the obesity level in the considered individuals. Each implemented model will be explained in the following pages and, finally, I will provide a summarized table for all the considered models in order to comprehend which one performed better.

Before diving into more detailed explanations, I computed the correlation between the data in order to highlight the dummy variable trap, a problem that may arise when techniques such as one-hot encoding or ordinal encoding are applied to categorical variables. This relatively known problem turns out to be easily manageable, since one of the best approaches is to consider only $K - 1$ dummy variables when encoding a categorical variable with K response classes. Therefore, I proceed to remove the following dummy variables:

- "GenderMale";
- "family_history_with_overweightno";
- "FAVCno";
- "Smokeno";
- "SCCno";
- "MTRANSWalking", the only dummy variable excluded from the ordinal encoding step.

The Supervised Learning approaches that I implemented in order to make inference are:

- Linear Regression;
- Stepwise Linear Regression;
- LASSO Linear Regression;
- Decision Tree;
- Random Forest.

1.3 Linear Regression

1.3.1 Simple Linear Regression

At this point our data set has 20 features and 1.950 observations. First of all, we start by splitting the dataset into a training set, which is randomly generated by including 80% of the total observations, and a test set, which has the remaining 20% of the observations.

```
> dim(test_data)
[1] 389  20
> dim(train_data)
[1] 1561  20
```

Figure 5: Training set and Test set dimensions.

The first Linear Regression model that we run includes all the 19 features and it is trained on the train set.

Figure 7 to Figure 9 shows some of the diagnostic on the Linear Regression model. It is noticeable a collinearity problem, also confirmed in Figure 8.

As we can see from Figure 6, some of the features included in the Linear Regression model turns out to be statistically insignificant, that is, we may exclude them from our Linear Regression model in order to improve our fit. Moreover, we can see that approximately 96% of the variance is explained by the considered features.

By diving deeper into the generated fit we can notice that the main important features regards physical attributes, such as Height and Weight, but also the Age of the individuals. Surprisingly, the model did not consider various health habits to be statistically significant, although individual's Gender seems to have a little role in predicting obesity, with Males being more subject to this risk.

It seems that having previous or current cases of obesity in the family increases the risk of incurring into obesity problems, but also individual's age (i.e., the older it is, the higher the risk of obesity). Lastly, consuming more meals during the day may lead to higher risk.

In the next sections we will deal with the insignificant features by implementing Stepwise and LASSO approaches to the Linear Regression model.


```

Call:
lm(formula = NObeyesdad ~ ., data = train_data)

Residuals:
    Min       1Q   Median       3Q      Max
-2.15433 -0.27280  0.00656  0.26539  1.74405

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    8.4058397   0.3205138   26.226 < 2e-16 ***
GenderFemale   -0.0391868   0.0148161    -2.645  0.00826 **
Age             0.0297857   0.0031863     9.348 < 2e-16 ***
Height        -7.4044427   0.1877735   -39.433 < 2e-16 ***
Weight         0.0762834   0.0006236  122.334 < 2e-16 ***
family_history_with_overweightyes  0.1296325   0.0133242     9.729 < 2e-16 ***
FAVCyes        0.0073720   0.0118607     0.622  0.53433
FCVC           0.0046327   0.0119734     0.387  0.69887
NCP             0.0303537   0.0115722     2.623  0.00880 **
CAEC          -0.0659815   0.0119822    -5.507 4.28e-08 ***
SMOKEyes       -0.0168590   0.0119268    -1.414  0.15770
CH2O           0.0010716   0.0114044     0.094  0.92515
SCCyes        -0.0090180   0.0113328    -0.796  0.42631
FAF           -0.0788982   0.0121400    -6.499 1.09e-10 ***
TUE           -0.0106891   0.0113773    -0.940  0.34761
CALC          -0.0495735   0.0118756    -4.174 3.15e-05 ***
MTRANSAutomobile -0.0308286   0.0267811    -1.151  0.24986
MTRANSBike     -0.0032149   0.0113243    -0.284  0.77653
MTRANSMotorbike  0.0108023   0.0113201     0.954  0.34010
MTRANSPublic_Transportation  0.0514304   0.0268500     1.915  0.05562 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4303 on 1541 degrees of freedom
Multiple R-squared:  0.9558,    Adjusted R-squared:  0.9553
F-statistic: 1756 on 19 and 1541 DF,  p-value: < 2.2e-16

```

Figure 6: Linear Regression model on all of the 19 features, with response variable NObeyesdad.

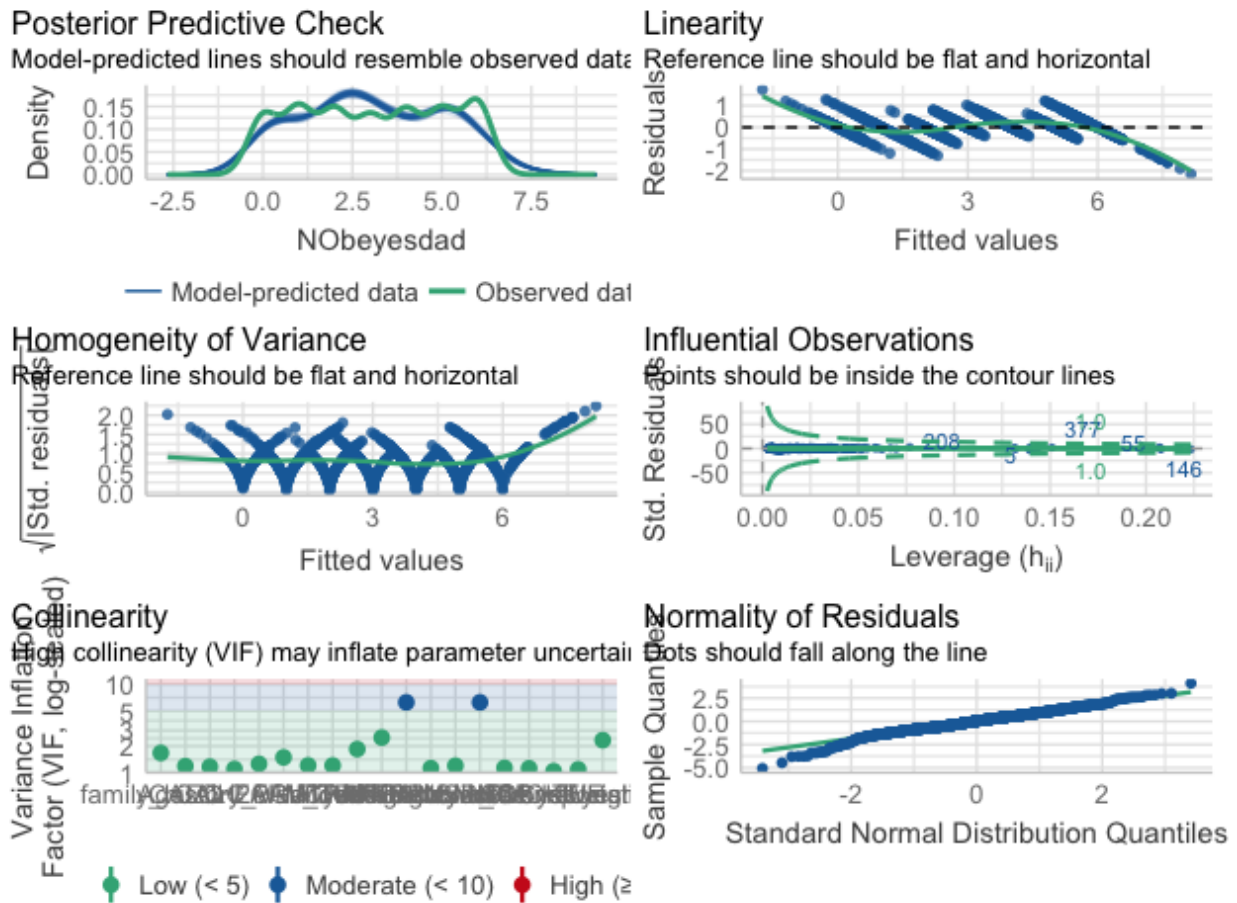


Figure 7: Linear Regression Model performance on different aspects.

```
> check_heteroscedasticity(lin_mod)
Warning: Heteroscedasticity (non-constant error variance) detected (p < .001).
> check_autocorrelation(lin_mod)
Warning: Autocorrelated residuals detected (p < .001).
```

Figure 8: Heteroscedasticity and Autocorrelation check on Linear Regression model.

Moderate Correlation

	Term	VIF	VIF 95% CI	Increased SE	Tolerance	Tolerance 95% CI
	MTRANSBike	1.12	[1.08, 1.20]	1.06	0.89	[0.83, 0.93]
	MTRANSPublic_Transportation	6.25	[5.72, 6.84]	2.50	0.16	[0.15, 0.17]

Figure 9: Moderate correlations resulting from Autocorrelation check on Linear Regression model.

1.3.2 Stepwise Regression

Stepwise linear regression is a method of fitting multiple regression variables while simultaneously removing the variables that results to be less explanatory. As I mentioned above, performing simple Linear Regression on all of the 19 considered features resulted in a regression in model in which approximately 10 features were insignificant. By implementing the stepwise approach, that is, a mixed technique that implement both backward and forward selection, I was able to achieve the same exact result as performing a simple Linear Regression model. As shown in Figure 10, in the resulted model almost all the selected features results to be effectively significant, with an exception for MTRANSMotorbike and SMOKEyes. Moreover, we can observe how the number of features decreased from 19 to 12.

```
> lin_step_sum

Call:
lm(formula = NObeyesdad ~ GenderFemale + Age + Height + Weight +
    family_history_with_overweightyes + NCP + CAEC + SMOKEyes +
    FAF + CALC + MTRANSMotorbike + MTRANSPublic_Transportation,
    data = train_data)

Residuals:
    Min       1Q   Median       3Q      Max
-2.16370 -0.27167  0.00343  0.26847  1.75450

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      8.3648859   0.3164537   26.433 < 2e-16 ***
GenderFemale     -0.0376138   0.0143121   -2.628  0.00867 **
Age               0.0298959   0.0030328    9.857 < 2e-16 ***
Height          -7.3897537   0.1862685  -39.673 < 2e-16 ***
Weight           0.0764353   0.0006002  127.340 < 2e-16 ***
family_history_with_overweightyes  0.1294937   0.0131444    9.852 < 2e-16 ***
NCP              0.0294123   0.0115088    2.556  0.01069 *
CAEC            -0.0661820   0.0117561   -5.630 2.14e-08 ***
SMOKEyes        -0.0182464   0.0117915   -1.547  0.12197
FAF             -0.0791130   0.0119141   -6.640 4.32e-11 ***
CALC            -0.0496976   0.0118033   -4.210 2.69e-05 ***
MTRANSMotorbike  0.0157269   0.0105327    1.493  0.13560
MTRANSPublic_Transportation  0.0797287   0.0127766    6.240 5.63e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4297 on 1548 degrees of freedom
Multiple R-squared:  0.9558,    Adjusted R-squared:  0.9554
F-statistic: 2786 on 12 and 1548 DF,  p-value: < 2.2e-16
```

Figure 10: Stepwise Linear Regression model on a selected subset of features, with response variable NObeyesdad.

As we can see from Figure 11 almost everything remains unchanged with respect to the simple Linear Re-

gression model. The only difference with the simple Linear Regression model is that the stepwise regression does not include any feature that have moderate or high level of correlation, as opposed to the simple Linear Regression model in which two variables were detected to be moderately correlated. Also the portion of explained variance is exactly the same as the one in the simple Linear Regression model, achieving an $R^2 \approx 96\%$.

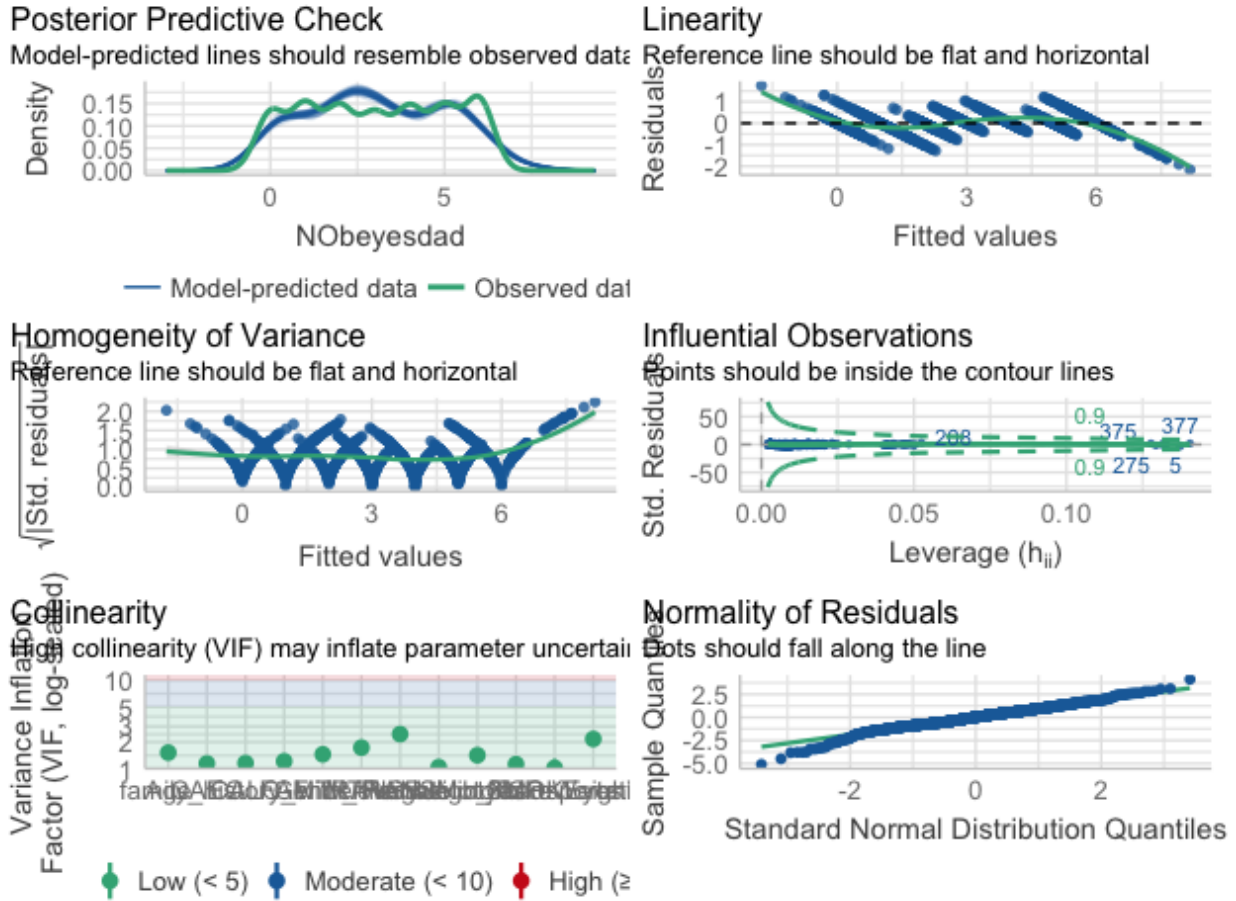


Figure 11: Stepwise Linear Regression model performance of different aspects.

1.3.3 LASSO Regression

LASSO, which stands for Least Absolute Shrinkage and Selection Operator, is a particular type of regression model in which shrinkage techniques are implemented. These techniques allow to shrink the coefficient estimates towards a central point (which is generally 0), and to obtain sparse models, that is, models in which we have fewer features.




Figure 12: Plot of the minimum $\log(\lambda)$ selected through cross validation on the training set

After we had computed the minimum value for the penalty term λ , the LASSO selected the features that are reported in Figure 13.

[1] "GenderFemale"	"Age"
[3] "Height"	"Weight"
[5] "family_history_with_overweightyes"	"FAVCyes"
[7] "NCP"	"CAEC"
[9] "SMOKEyes"	"SCCYes"
[11] "FAF"	"TUE"
[13] "CALC"	"MTRANSAutomobile"
[15] "MTRANSMotorbike"	"MTRANSPublic_Transportation"

Figure 13: Selected features using LASSO approach.

It is easy to see that in this case LASSO selects more features than the stepwise approach. Figure 14 does not highlight any major difference with respect to the simple Linear Regression model. Furthermore, Figure 15 highlights two regressors with moderate correlation, the same ones that was captured in the simple Linear Regression model.

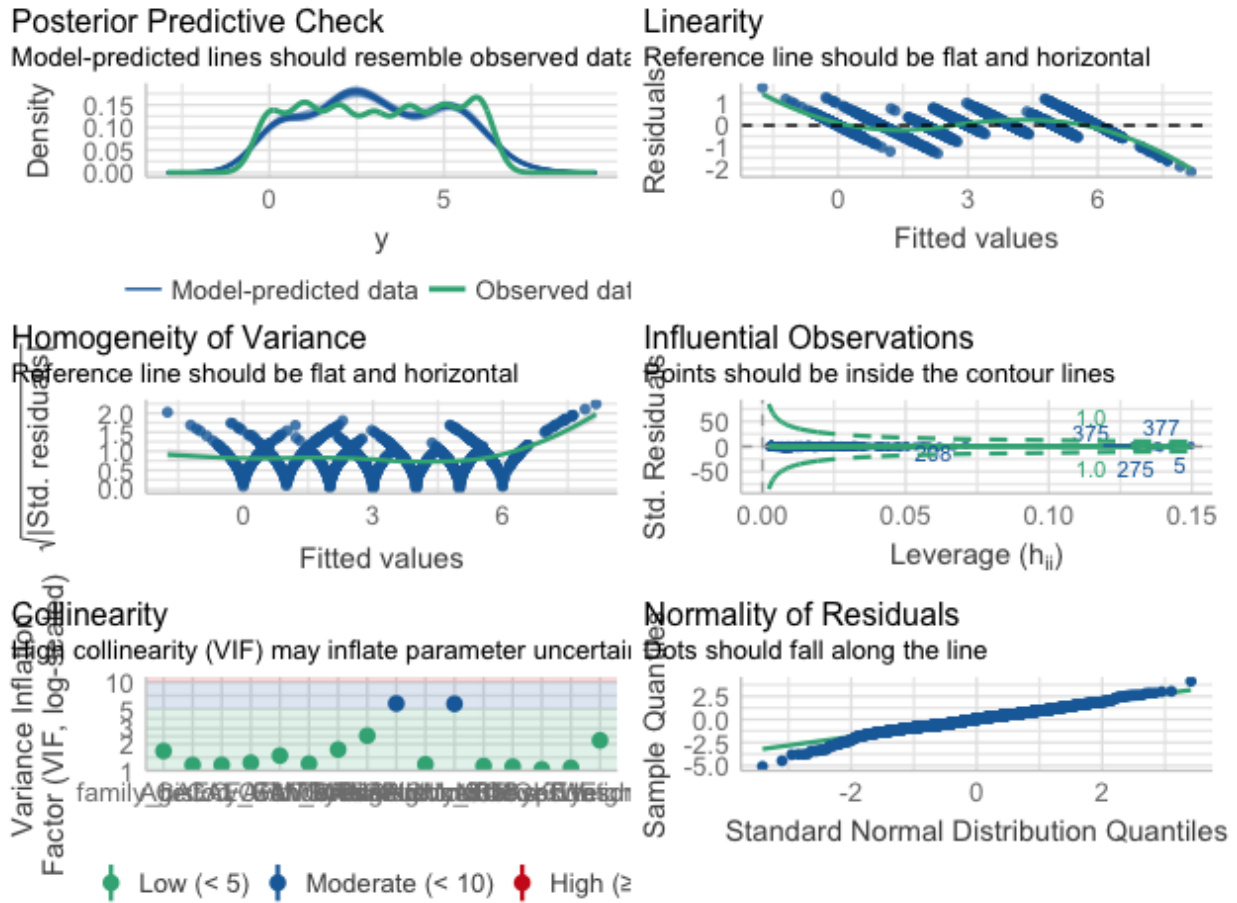


Figure 14: LASSO Linear Regression model performance of different aspects.

Moderate Correlation

Term	VIF	VIF 95% CI	Increased SE	Tolerance	Tolerance 95% CI
MTRANSMotorbike	1.18	[1.12, 1.26]	1.09	0.85	[0.79, 0.89]
MTRANSPublic_Transportation	5.72	[5.23, 6.26]	2.39	0.17	[0.16, 0.19]

Figure 15: Moderate correlations resulting from Autocorrelation check on LASSO Linear Regression model .

Finally, Figure 16 shows the LASSO regression model coefficients with the relative p -values. In this case no variables were excluded with respect to the simple Linear Regression model. Moreover, it seems that no big difference had occurred at all, even in the R^2 being approximately equal to 96%.

```

Call:
lm(formula = fmla, data = train_data)

Residuals:
    Min       1Q   Median       3Q      Max
-2.15513 -0.27285  0.00809  0.26682  1.74999

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    8.3986555   0.3196590   26.274 < 2e-16 ***
GenderFemale   -0.0375599   0.0143338    -2.620  0.00887 **
Age             0.0297837   0.0031793     9.368 < 2e-16 ***
Height        -7.4031530   0.1873537  -39.514 < 2e-16 ***
Weight         0.0763414   0.0006052  126.139 < 2e-16 ***
family_history_with_overweightyes  0.1294656   0.0132694     9.757 < 2e-16 ***
FAVCyes        0.0074027   0.0118160     0.627  0.53108
NCP             0.0302945   0.0115483     2.623  0.00879 **
CAEC          -0.0656003   0.0118337    -5.544 3.48e-08 ***
SMOKEyes       -0.0167839   0.0118839    -1.412  0.15806
SCCyes         -0.0086872   0.0112730    -0.771  0.44105
FAF            -0.0785613   0.0119945    -6.550 7.82e-11 ***
TUE            -0.0110751   0.0113138    -0.979  0.32778
CALC          -0.0495580   0.0118385    -4.186 3.00e-05 ***
MTRANSAutomobile -0.0288208   0.0256868    -1.122  0.26203
MTRANSMotorbike  0.0112571   0.0112225     1.003  0.31598
MTRANSPublic_Transportation  0.0535751   0.0256565     2.088  0.03695 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4299 on 1544 degrees of freedom
Multiple R-squared:  0.9558,    Adjusted R-squared:  0.9554
F-statistic: 2089 on 16 and 1544 DF,  p-value: < 2.2e-16

```

Figure 16: LASSO Regression model on a selected subset of features, with response variable NObeyedad.

1.3.4 Linear Regression models comparison

We have implemented a simple Linear Regression model in order to analyse and predict the obesity level in our data. Then we adapted our linear model by implementing two different techniques that aimed at reshaping our base model by selecting only the most important regressors or by shrinking those features that does not improve our fit.

At this point we are interested in comparing those three different models, in order to comprehend which approach leads to better results.

Comparison of Model Performance Indices

Name	Model	R ²	R ² (adj.)	RMSE	Sigma	AIC weights	BIC weights	Performance-Score
lin_mod_step	lm	0.956	0.955	0.428	0.430	0.922	1.000	66.67%
lin_lasso	lm	0.956	0.955	0.428	0.430	0.074	1.80e-06	56.08%
lin_mod	lm	0.956	0.955	0.428	0.430	0.004	3.31e-11	33.33%

Figure 17: Simple Linear Regression, LASSO Linear Regression and Stepwise Linear Regression comparison.

In Figure 17 are reported the various statistic for each of the three considered models. As we have already stated various time, the three models have the same R^2 . Also the RMSE turns out to be the same for the three models. The only noticeable differences are in the AIC weights³, which represents the model's relative probability of leading good results. Moreover, Stepwise Linear Regression approach register the highest "Performance Score"⁴, almost twice as the simple Linear Regression model. This could be attributed to the fact that the Stepwise approach selected only the effectively significant variables, excluding those that did not have a high explanatory power. Also the LASSO Linear Regression registered a good "Performance Score" with respect to the baseline, even though 16 of the 19 features were included.

Finally, Figure 18 allows us to visualize the results of the comparison using a Spider plot.

Comparison of Model Indices




Figure 18: Simple Linear Regression, LASSO Linear Regression and Stepwise Linear Regression comparison using a Spider web plot.

³<https://link.springer.com/content/pdf/10.3758/BF03206482.pdf>

⁴This score ranges from 0 to 100%. Note that all score value do not necessarily sum up to 100. Rather, calculation is based on normalizing all indices (i.e. rescaling them to a range from 0 to 1), and taking the mean value of all indices for each model. This is a rather quick heuristic, but might be helpful as exploratory index.

1.4 Classification

In statistics, the logistic model is a statistical tool that models the probability of one event taking place by having the log-odds for the event to be a linear combination of one or more independent variables. On the other hand, Logistic Regression consists in estimating the parameters of a logistic model.

In this case, the response feature NObeyesdad is a categorical variable, that is, it has a discrete number of possible outcomes. Moreover, in this particular case the response features does not happen to be a binary response variable (i.e. it only has $K = 2$ response classes), but rather it has $K = 7$ different response classes. I will briefly dive into a specific type of classification algorithm, named Multinomial Logistic Regression.

The main difference with respect to the binary Logistic Regression is that Multinomial Logistic Regression allows to model the probability of an event that has multiclass outcomes ($K > 2$). In order to implement this model I used the *nnet*⁵ package from Rstudio.

Before diving any further into our Multinomial Logistic Regression model, I decided to duplicate the previously created training and test set. This will be helpful later on, since we will develop the model without encoding and factorizing our response variable NObeyesdad, utilizing the function *relevel()*⁶ to assign different values to our classes in the response variable.

```
multinom_model = multinom(NObeyesdad ~ ., data = train_data_log)
```

Figure 19: Multinomial Logistic Regression call in Rstudio.

Figure 19 shows the Multinomial Logistic Regression model call in Rstudio. Once we trained our model, we applied the exponential function to the coefficients in order to obtain the relative odds. The resulting model applied on the training set turns out to have an accuracy⁷ of 96.82% on the train data.

R code 1.1: Trained Multinomial Logistic Regression outcome on train and test data.

```
1 > tab = table(train_data_log$NObeyesdad, train_data_log$ClassPredicted)
2 > tab
3
4           Insufficient_Weight Normal_Weight Overweight_Level_I
5 Insufficient_Weight           225           6           0
6 Normal_Weight                7          207           8
7 Overweight_Level_I            0           11          183
8 Overweight_Level_II           0           0           2
9 Obesity_Type_I               216           0           0
10 Obesity_Type_II              3           0           0
11 Obesity_Type_III             0           0           0
12
13           Obesity_Type_I Obesity_Type_II
14 Insufficient_Weight      0           0
15 Normal_Weight           0           0
16 Overweight_Level_I       0           0
17 Overweight_Level_II      2           0
```

⁵<https://cran.r-project.org/web/packages/nnet/nnet.pdf>

⁶The levels of a factor are re-ordered so that the level specified by ref is first and the others are moved down. This is useful for *contr.treatment* contrasts which take the first level as the reference.

⁷Sum of the diagonal observations divided by the sum of the total observations.

```

18   Obesity_Type_I           218           0
19   Obesity_Type_II          0          259
20   Obesity_Type_III         0           0
21 > round((sum(diag(tab))/sum(tab))*100,2)
22 [1] 96.82
23
24 > tab_1 = table(test_data_log$NObeyesdad, test_data_log$ClassPredicted)# Building
      classification table
25 > tab_1
26
27           Insufficient_Weight Normal_Weight Overweight_Level_I
      Overweight_Level_II Obesity_Type_I
28 Insufficient_Weight           46           3           0
29           0           0
30 Normal_Weight           1          45           2
31           0           0
32 Overweight_Level_I           0           6          47
33           3           0
34 Overweight_Level_II           0           0           4
35           56           3
36 Obesity_Type_I           0           0           0
37           1          45
38 Obesity_Type_II           0           0           0
39           2           0
40 Obesity_Type_III           0           0           0
41           0           0
42
43           Obesity_Type_II
44 Insufficient_Weight           0
45 Normal_Weight           1
46 Overweight_Level_I           0
47 Overweight_Level_II           0
48 Obesity_Type_I           1
49 Obesity_Type_II          62
50 Obesity_Type_III           0
51 > round((sum(diag(tab_1))/sum(tab_1))*100,2)
52 [1] 91.77

```

Then, in order to test our trained model, we applied the Multinomial Logistic Regression to the test data. The resulting accuracy, computed as before, is 91.77%.

1.5 Beyond Linearity

Now I will move to a different class of Supervised Learning techniques. This kind of approaches have the same basis as the ones presented until this point, except from the fact that they allow to consider non-linear effects on the model. In particular, this kind of model has the ability to intercept different interactions between features that a simple linear model is not able to consider, therefore improving the interpretability of the model. This advantage comes at a cost, that is, the increase in complexity.

1.5.1 Decision Trees

Decision tree is a Supervised Learning technique used as a predictive model to draw inference from a set of observations.

Decision tree builds regression or classification models in the form of a tree structure. It breaks down a dataset into continously smaller subsets while at the same time an associated decision tree is incrementally developed. The final result is a tree with decision nodes, which represents values for the attribute tested, and leaf nodes, which represent a decision on the numerical target.

In this case I did not use the *tree* package in Rstudio, but rather the *rpart*⁸ package.

⁸<https://cran.r-project.org/web/packages/rpart/rpart.pdf>.

```

1 > printcp(tree_1)
2
3 Classification tree:
4 rpart(formula = NObeyesdad ~ ., data = train_data_log, method = "class")
5
6 Variables actually used in tree construction:
7 [1] Age          CAEC          FAVCyes          FCVC          GenderFemale Height
8      Weight
9
10 Root node error: 1092/1351 = 0.80829
11
12 n=1351 (210 osservazioni eliminate a causa di valori mancanti)
13
14      CP nsplit rel error  xerror    xstd
15 1  0.211538    0  1.00000 1.00000 0.013250
16 2  0.198718    1  0.78846 0.78938 0.016176
17 3  0.132784    2  0.58974 0.60440 0.016825
18 4  0.064103    3  0.45696 0.45971 0.016265
19 5  0.036630    4  0.39286 0.41575 0.015899
20 6  0.027015    7  0.28297 0.26832 0.013872
21 7  0.020147    9  0.22894 0.22070 0.012886
22 8  0.016484   11  0.18864 0.19963 0.012382
23 9  0.014652   12  0.17216 0.19231 0.012196
24 10 0.010989   13  0.15751 0.18223 0.011929
25 11 0.010531   14  0.14652 0.17491 0.011727
26 12 0.010073   16  0.12546 0.17125 0.011624
27 13 0.010000   17  0.11538 0.15476 0.011135

```

As we can see in the above piece of code, the Classification tree selects 7 features in order to construct the tree, that is, these features will pop-up in the upper part of the tree and they will resemble the most important decisions splits. Also, Figure 20 reports a brief summary of the computed tree where can see that the Complexity Parameter⁹ *CP* decreases significantly when the size of the tree is equal to 10, and then slowly approaches its minimum as the size of the tree reaches its maximum.

⁹The complexity parameter (*cp*) in *rpart* is the minimum improvement in the model needed at each node.




Figure 20: Complexity parameter for the Classification Tree.

We then proceeded to prune the trained tree at a complexity level $cp = 0.020$, that is at the point in which increasing the complexity parameter does not lead any significant increase in the tree's performance.

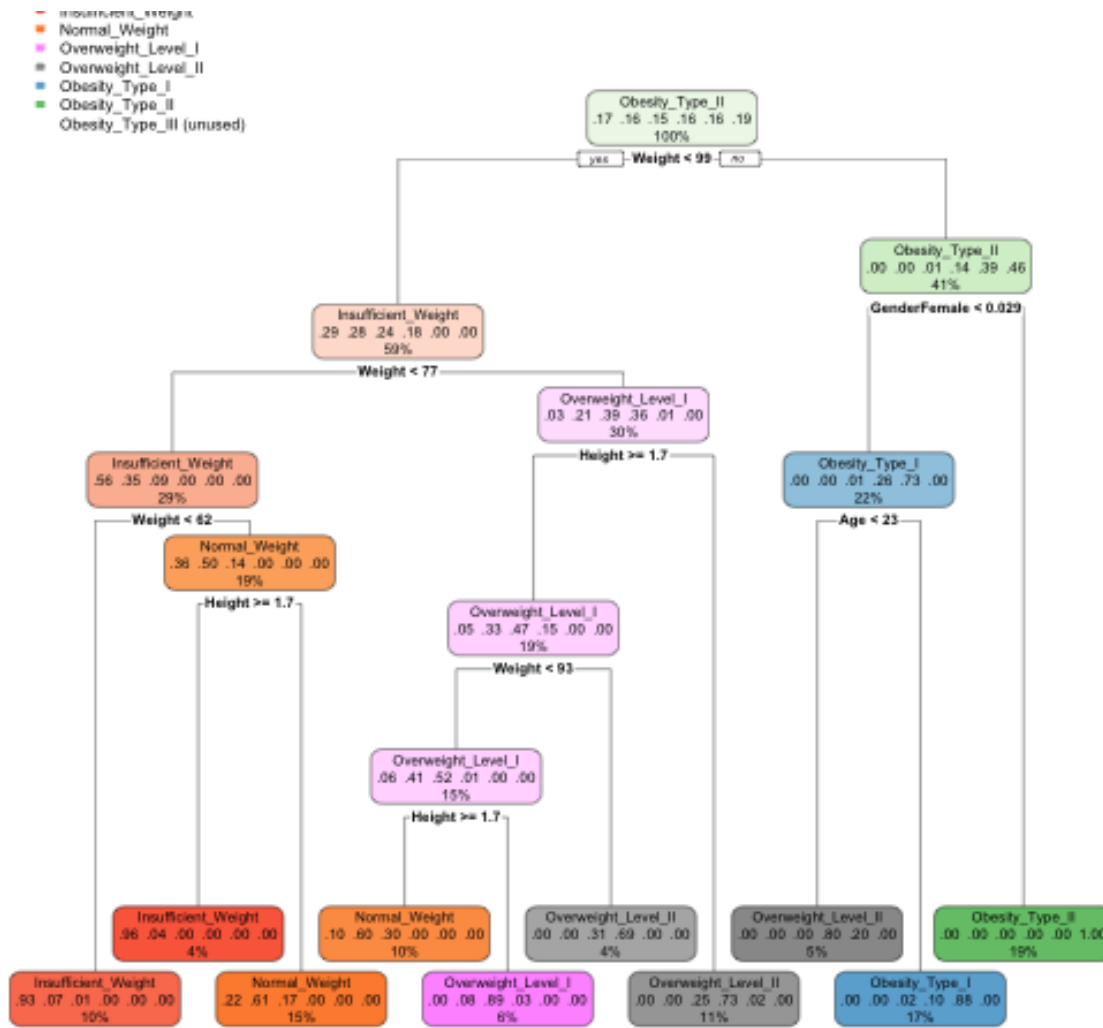


Figure 21: Pruned Classification Tree with $cp = 0.020$.

As we can see from Figure 21, with the selected level of Complexity we obtain a pruned tree with just 10 terminal nodes. Moreover, it is noticeable that the "Obesity_Type.III" response class is not even considered in the pruned tree. Also, the features that are used in the decision nodes drops from 7 to just 4, with Weight being the root of the tree. The most important features turns out to be Weight and Height, but this should be obvious since the BMI index relies on both of these features in order to be computed. Surprisingly, high level of BMI corresponding to Obesity Tipe II results to affect mostly males than females.

R code 1.3: RMSE, MAE and R^2 .

```

1 > TREE = c(rmse(test_data$NObeyesdad, dt_1_pred),
2 +         mae(test_data$NObeyesdad, dt_1_pred),
3 +         R2(train_data$NObeyesdad, dt_1_yhat))
4 > tab_2 = table(test_data_log$NObeyesdad,dt_1_pred)
5 > TREE
6 [1] 0.5554127 0.2827763 0.9578041
7 > tab_2 = table(test_data_log$NObeyesdad,dt_1_pred)
8 > tab_accuracy = sum(diag(tab_2))/sum(tab_2)
9 > tab_accuracy
10 [1] 0.8658537

```

Finally, we can see from the above piece of code the Root Means Squared Error RMSE¹⁰, the Mean Absolute Average MAE¹¹ and the R^2 . I also computed the accuracy computing the MSE on the predicted observations on the test data with respect to the actual class value. The train Decision Tree performs quite well, with an accuracy level of approximately 87%.

1.5.2 Random Forest

Random forests algorithm is an ensemble learning method that operates by constructing a multitude of decision trees at training time. For regression tasks, the mean or average prediction of the individual trees is returned.

As always, the Random Forest is trained on the training data. At each split of tree it is possible to select a number of features equal to the square root of the total number of features, being $\sqrt{19} = 4$.

R code 1.4: Random Forest

```

1 > rf_1 = randomForest(NObyesdad ~ ., data=train_data, mtry = sqrt(19))
2 > rf_1
3
4 Call:
5 randomForest(formula = NObyesdad ~ ., data = train_data, mtry = sqrt(19))
6           Type of random forest: regression
7           Number of trees: 500
8 No. of variables tried at each split: 4
9
10           Mean of squared residuals: 0.09918575
11           % Var explained: 97.6

```




Figure 22: Random Forest Error with respect to the Number of Trees.

¹⁰RMSE is the standard deviation of the residuals (prediction errors). Residuals are a measure of how far from the regression line data points are; RMSE is a measure of how spread out these residuals are. Since the errors are squared before they are averaged, the RMSE gives a relatively high weight to large errors. This means the RMSE should be more useful when large errors are particularly undesirable.

¹¹MAE measures the average magnitude of the errors in a set of predictions, without considering their direction. It's the average over the test sample of the absolute differences between prediction and actual observation where all individual differences have equal weight.

We can see that the Number of trees turns out to be 500 but, as stated in Figure 22, it is more than sufficient to minimize the Random Forest error. In fact, with at least 100 trees generated by the Random Forest algorithm we have a significant decrease in the error.

As opposed to Decision Trees algorithm where the model is also supported by a graphical representation that allows to facilitate the readability and interpretability of the model, Random Forest does not provide such a feature. In order to comprehend which regressors are significant we can rely on Figure 23, which represents all the features included in the Random Forest and their relative importance, named *IncNodePurity*. Once again the most important feature by far is "Weight". It turns out that cases of obesity in the family may be significant in order to predict future obesity problem in the individual. This feature was not considered at all by the Decision Tree algorithm, but in all of the three Linear Regression models was considered to be statistically significant.




Figure 23: Random Forest include variables and relative importance.

Then, we applied the trained Random Forest algorithm to the test data in order to understand how it performs on new observations. Figure 24 displays the **Out-Of-Bag Error**, that is, the prediction error of the Random Forest, and also the true value of the response variable in the test set. As we can see, the model performs quite well.




Figure 24: Random Forest OOB and Test Error.

1.6 Conclusions

In this first part of the report I implemented different Supervised Learning techniques in order to design different models that aims at predicting a possible Obesity problem in an individual. First of all, I implemented a simple Linear Regression model and, in order to improve the interpretability of the findings, I considered Stepwise and LASSO approaches to the Linear Regression model. The conclusions about these three models are reported in the previous sections.

Then, I considered a different class of methods, mainly known for capturing the non-linear relation between different features. The two selected models are Decision Tree and Random Forest. These two models allowed us to better capture any non-linear relationship between features. This resulted in various features that were considered by these class of methods and that did not appear to be significant in the linear model. Moreover, the main features selected to be the more significant across all the five models implemented turns out to be the same: Height, Weight and Age of the individuals plays a major role in predicting a possible Obesity problem. This should not be surprising since, as stated various time before, the Body Mass Index is directly computed from two of these three features. Also, Age plays an important role in predicting obesity level since it is a known fact that as an human gets old, it starts to lower his physical activities.

An interesting fact is that various features that resembles healths habits such as smoking, or movement transportations means such as preferring to walk instead of driving, did not played a major role in our models. Also different eating habits, such as consuming vegeatables during meals or measuring the consumed calories, played a little if not any role.

Considering the performance of the implemented models, all the five models had an overall high explanatory power and did performe pretty well in predicting new observations. Moreover, Random Forest turns out to

be the best model by all of the three performance parameters that I reported. It has by far the lowest RMSE and MAE, and the R^2 is little higher than the one reported for the Linear Regression models. Overall, in terms of accuracy power on unobserved observations, Decision Tree performed worst than all of the other models, striking a 55.5% RMSE.

R code 1.5: Model Comparisons

```
1 > metrics
2           OLS   Stepwise    LASSO    Robust      TREE      RF
3 RMSE 0.4620516 0.4621797 0.4616688 0.4678218 0.5554127 0.3673668
4 MAE  0.3474481 0.3481118 0.3472340 0.3431884 0.2827763 0.2174774
5 R2    0.9558427 0.9557519 0.9558357 0.9555305 0.9578041 0.9689971
```

2 Part 2: Unsupervised Learning

Unsupervised Learning represents a variety of different algorithm that aims at learning patterns from the unlabeled data in order to obtain meaningful insights.

Considering our case, the response variable NObeyesdad resulted to be extremely useful when we applied Supervised Learning techniques to our data. Now, for this second part of the report I will try to perform Unsupervised Learning technique in order to cluster the data without relying on the information provided by the NObeyesdad feature which, as stated in the previous part, is directly computed from an individual's Height and Weight measurement.

The Unsupervised Learning algorithms that I will use in this second part are:

- **MCA: Multiple Correspondence Analysis;**
- **Hierarchical Clustering.**

Before starting with the model construction and with the analysis, I decided to reload the original data without any type of transformation. Then, I substituted some of the numeric feature values with the corresponding meaning, that is, I transformed the majority of the features in qualitative ones¹². The only features that I did not change are: Height, Age, Weight, and those who had a nominal binary response (i.e. "yes" or "no").

2.1 Multiple Correspondence Analysis

Multiple Correspondence Analysis is a data analysis technique for nominal categorical data used to detect and represent underlying structures in a data set. It can be seen as the counter part of the Principal Component Analysis, a widely known technique used in data analysis to reduce the dimensionality of the data. Moreover, this latter approach is used mainly on continuous variables.

Multiple Correspondance Analysis is a variation of the Principal Component Analysis that could be applied to nominal qualitative features, such as in our case. In order to implement this model I used the *FactoMineR*¹³ package in Rstudio.

In this part I decided to use as *active variables*, that is the ones used in the MCA, the one related to eating behavior, while the remaining ones are assigned to be *supplementary information*, that is variables that will not be used directly in the model, but for which coordinates will be predicted. Therefore, the active variables are: FAVC, FCVC, NCP, CAEC, CH2O, SCC, and CALC.

¹²Technically Multiple Correspondence Analysis is obtained by using a standard correspondence analysis on an indicator matrix (i.e., a matrix whose entries are 0 or 1). It is used to analyze a set of observations described by a set of nominal variables. Each nominal variable comprises several levels, and each of these levels is coded as a binary variable. For example gender, (F vs. M) is one nominal variable with two levels. The pattern for a male respondent will be 0 1 and 1 0 for a female. The complete data table is composed of binary columns with one and only one column taking the value "1" per nominal variable. See <https://personal.utdallas.edu/~herve/Abdi-MCA2007-pretty.pdf> for more details.

¹³<https://cran.r-project.org/web/packages/FactoMineR/FactoMineR.pdf>.

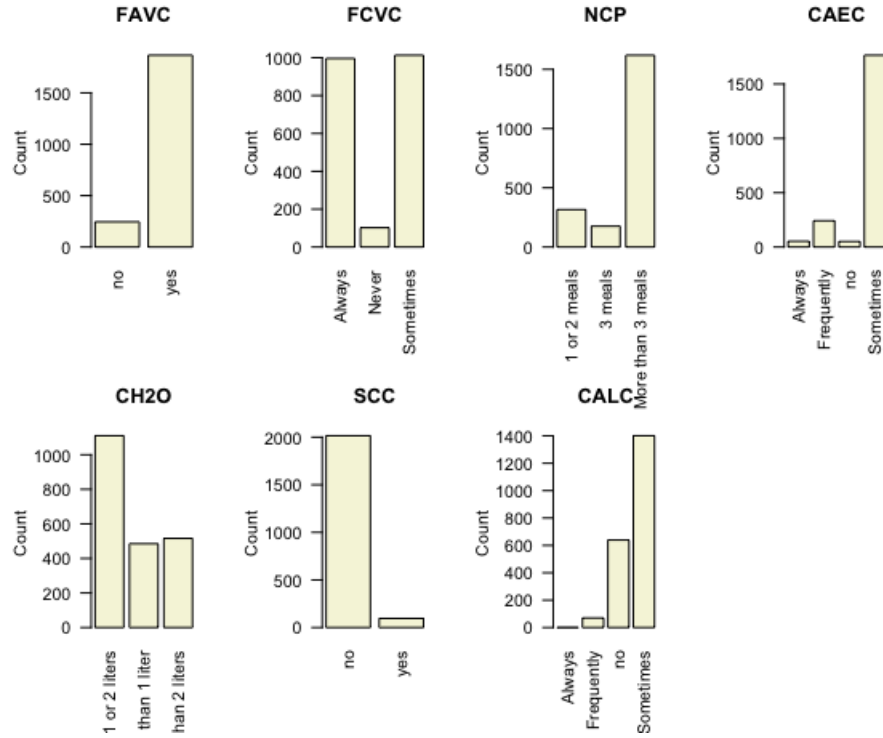


Figure 25: Frequency for the *Active Variables*.

I proceeded to run the model using the *MCA* function from the *FactoMineR* package which automatically detects the number of categorical variables (7 in our case). Since MCA (but also PCA) are known for being very sensitive to large outliers, I used the *outliers* function from the *FactoInvestigate*¹⁴ package to detect them, that is, detect singular individuals that concentrates too much variance, properly called inertia. It resulted that 101 individuals are considered to be outliers¹⁵, therefore I proceeded to remove them from the data.

Figure 26 (left) displays the percentage of explained variance in each dimension in the model, represented by the correspondent eigenvalue. Clearly, the first two dimensions are capable to explain approximately the 22% of the comprehensive variance in the data, measured by the percentage of variance explained in the model. In order to obtain a "good" Multiple Correspondence Analysis, we must consider using at least more than this two dimensions, since the majority of the variance is still unexplained. Moreover, considering at least ten dimensions will lead to explain the 87% of the variance in our data. This apparently bad performance can be explained by the fact that our original data have a small number of feature (17 in the beginning), and before applying the MCA algorithm we selected a subset with 7 active variables. Starting with these low number of features it is quite difficult to achieve a good performance in terms of dimension reduction without reducing the explanation of the variance.

R code 2.1: MCA summary, eigenvalues

```

1 > summary(res.mca)
2
3 Call:
4 MCA(X = data_active, ncp = 3, graph = TRUE)
5
6
```

¹⁴<https://cran.r-project.org/web/packages/FactoInvestigate/FactoInvestigate.pdf>.

¹⁵The aforementioned function plots also the main features for which each individual results to concentrate inertia. The obtained results are not shown in the report but the function call is available in the Appendix code.

7	Eigenvalues							
8		Dim.1	Dim.2	Dim.3	Dim.4	Dim.5	Dim.6	Dim.7
9	Variance	0.197	0.171	0.164	0.159	0.147	0.144	0.138
10	% of var.	11.512	9.966	9.554	9.252	8.574	8.410	8.032
11	Cumulative % of var.	11.512	21.479	31.033	40.285	48.859	57.269	65.301
12		Dim.8	Dim.9	Dim.10	Dim.11	Dim.12		
13	Variance	0.132	0.124	0.120	0.118	0.101		
14	% of var.	7.706	7.241	6.984	6.879	5.889		
15	Cumulative % of var.	73.007	80.248	87.232	94.111	100.000		

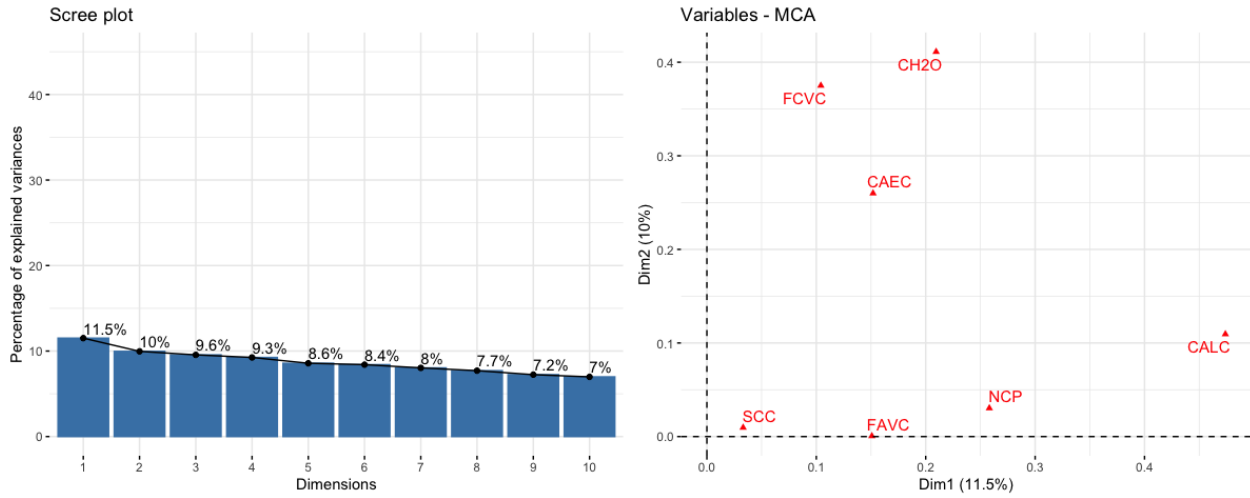


Figure 26: *Left*: Eigenvalues-Variance for MSE in each dimension. *Right*: Correlation between variables and principal dimensions.

On the other hand, Figure 26 (right) shows the correlation between the active variables and the first two dimensions. As we can see, the FCVC, CH2O and CAEC seems to be more correlated with the second dimension, while CALC and NCP with the first one. We can consider that good eating habits are better captured by the second dimension, while other aspects such as alcohol consumption or the number of meals per day are more concentrated in the first dimension. FAVC and also SCC seems to be correlated with both of the dimensions, that is, it may that higher dimension capture them the best way.

Figure 27 reports the quality of representation of variable categories, computed from the *squared cosine* (\cos^2), which measures the degree of association between variable categories and a particular axis (dimension). Below the figure there is a detail about each variable category with the related dimension. FAVC impacts only the first dimension and this explains the fact that is closest to the x-axis; also we can see that it reports 0 in the second dimension and 0.08, confirming the previous statement. It is noticeable that daily alcohol consumption seems to not influence the first three dimensions, while no alcohol or sporadic consumption shows their effect on the first dimension.




Figure 27: Quality of representation of variable categories.

R code 2.2: Model Comparisons

```

1 > round(var$cos2,2)
2
3 Dim 1 Dim 2 Dim 3
4 FAVC_no 0.15 0.00 0.08
5 FAVC_yes 0.15 0.00 0.08
6 FCVC_Always 0.10 0.34 0.06
7 FCVC_Never 0.00 0.10 0.03
8 FCVC_Sometimes 0.10 0.20 0.10
9 1_or_2_meals 0.20 0.03 0.00
10 3_meals 0.03 0.01 0.34
11 More_than_3_meals 0.24 0.01 0.15
12 CAEC_Always 0.02 0.12 0.25
13 CAEC_Frequently 0.13 0.15 0.18
14 CAEC_Sometimes 0.15 0.04 0.39
15 1_or_2_liters 0.04 0.39 0.00
16 Less_than_1_liter 0.05 0.24 0.00
17 More_than_2_liters 0.20 0.06 0.00
18 SCC_no 0.03 0.01 0.01
19 SCC_yes 0.03 0.01 0.01
20 CALC_Frequently 0.01 0.09 0.13
21 CALC_no 0.44 0.03 0.04
22 CALC_Sometimes 0.47 0.00 0.00

```

Looking at the individuals, we can see in Figure 28 the contribution, once again measured using squared cosine. The plot displays the individuals contribution, grouping the observation on the variable FAVC. As we could imagine, the observations are distributed along the x-axis. Moreover, people who tend to not consume frequently high caloric meals are positively correlated with the first dimension.




Figure 28: Quality of representation of individuals.

2.2 Hierarchical Clustering on Principal Components

Hierarchical Clustering is an algorithm that groups similar objects into groups called clusters. In particular clusters are considered to have different aspects one from the other, and each object contained in a certain cluster have more common things across other objects in the same cluster than objects contained in different clusters.

By treating each observation as a single cluster, the algorithm tries to find the two clusters that are closest together, in order to merge them in a single cluster. This operation is iterated until no merge is possible between two clusters.

In order to perform Hierarchical Clustering on Principal Components, we will continue the previous work obtained applying Multiple Correspondence Analysis.




Figure 29: Dendrogram obtained after Hierarchical Clustering.

Figure 29 displays the coloured dendrogram obtained from the Multiple Correspondence Analysis. Compared to other algorithms, such as K -means, Hierarchical Clustering does not require to specify the cluster number. The algorithm detects 4 different clusters, each represented with a different colour. At Height being around 0.10, we can see that two main clusters are detected. Figure 30 displays the factor map for each considered individual.




Figure 30: Factor map for the Hierarchical Clustering.

As we can see from the chunk of code below, clusters can be well described by the variables and/or categories, by the principal axes and by individuals. The first cluster appear to be composed by individuals who consume more than 3 meals on the daily basis, some of them being high-caloric meals, and whom consume alcohol in some occasions. On the other hand, the second clusters groups individuals who does not drink alcohol and who prefer to consume one or two meals per day. This consideration are assumed on the basis of the relevancy for each category in the cluster. In fact, we can see that in the first cluster are grouped individuals who drinks more than two liters per day of water, but also those that consume's less than one liter per day. In other words, the obtained clusters overlaps because different individuals with different habits turns out to be included in the same cluster. This happens because largest source of variation is similar among different groups.

These common variation seems to be concentrated in the second cluster, that is the one who overlaps with all the other three clusters. In fact, if we imagine for a bit to exclude it from Figure 30, we can clearly see that the three clusters are quite definite and not overlapping, exception made for a bunch of individuals that are considered to be in the fourth and in the third cluster.

Lastly, in the below chunk of code we can see some details about some individuals in each cluster. In particular we can look at the individuals who results to be closest to the centre of each cluster (indicated by *para*) and those who are more distant from the centre (via *dist*), with the relative distance.

R code 2.3: HC

```

1 > #Description of the 10 most relevant variables
2 > head(res.HCPC$desc.var$test.chi2,10)
3       p.value df
4 CALC  0.000000e+00 6
5 CAEC  2.380674e-307 6
6 FCVC  5.479560e-135 6
7 CH2O  9.911078e-132 6
8 NCP   9.060267e-119 6
9 FAVC  1.318048e-40 3
10 SCC  3.167638e-12 3
11
12 > #Description of the 10 most relevant categories
13 > head(res.HCPC$desc.var$category,10)

```


14	\$ '1'					
15		Cla/Mod	Mod/Cla	Global	p.value	v.test
16	CALC=CALC_Sometimes	57.379518	98.8326848	66.333666	8.403490e-167	27.526743
17	FCVC=FCVC_Always	61.684211	76.0051881	47.452547	2.300822e-94	20.608537
18	NCP=More_than_3_meals	47.840103	96.2386511	77.472527	3.149713e-68	17.455074
19	CH20=More_than_2_liters	67.446809	41.1154345	23.476523	3.127435e-48	14.592664
20	CAEC=CAEC_Sometimes	43.238636	98.7029831	87.912088	3.660482e-40	13.265688
21	FAVC=FAVC_yes	40.818584	95.7198444	90.309690	1.062135e-11	6.797821
22	SCC=SCC_no	39.125964	98.7029831	97.152847	6.006064e-04	3.431340
23	SCC=SCC_yes	17.543860	1.2970169	2.847153	6.006064e-04	-3.431340
24	FCVC=FCVC_Never	18.390805	2.0752270	4.345654	4.031692e-05	-4.105656
25	CAEC=CAEC_Always	0.000000	0.0000000	2.347652	8.382741e-11	-6.493567
26	FAVC=FAVC_no	17.010309	4.2801556	9.690310	1.062135e-11	-6.797821
27	CALC=CALC_Frequently	0.000000	0.0000000	2.947053	1.996402e-13	-7.349037
28	CAEC=CAEC_Frequently	5.128205	1.2970169	9.740260	9.404716e-30	-11.329212
29	NCP=1_or_2_meals	9.352518	3.3722438	13.886114	1.289401e-31	-11.699016
30	NCP=3_meals	1.734104	0.3891051	8.641359	6.393022e-34	-12.141139
31	CH20=1_or_2_liters	25.208526	35.2788586	53.896104	2.885254e-40	-13.283515
32	FCVC=FCVC_Sometimes	17.512953	21.9195850	48.201798	1.062408e-80	-19.024825
33	CALC=CALC_no	1.463415	1.1673152	30.719281	4.742403e-146	-25.735024
34						
35	\$ '2'					
36		Cla/Mod	Mod/Cla	Global	p.value	v.test
37	FCVC=FCVC_Sometimes	53.886010	80.996885	48.201798	4.854305e-95	20.683726
38	CH20=1_or_2_liters	51.065802	85.825545	53.896104	1.345431e-93	20.522866
39	CAEC=CAEC_Sometimes	36.477273	100.000000	87.912088	1.098423e-44	14.024851
40	NCP=3_meals	76.878613	20.716511	8.641359	7.107568e-37	12.685591
41	FAVC=FAVC_yes	34.292035	96.573209	90.309690	2.185631e-12	7.022097
42	FCVC=FCVC_Never	63.218391	8.566978	4.345654	1.288682e-09	6.068802
43	CALC=CALC_no	38.699187	37.071651	30.719281	2.735689e-05	4.194422
44	SCC=SCC_no	32.647815	98.909657	97.152847	5.568374e-04	3.451811
45	CALC=CALC_Sometimes	29.894578	61.838006	66.333666	3.637718e-03	-2.907980
46	NCP=More_than_3_meals	30.109607	72.741433	77.472527	5.808125e-04	-3.440420
47	SCC=SCC_yes	12.280702	1.090343	2.847153	5.568374e-04	-3.451811
48	CALC=CALC_Frequently	11.864407	1.090343	2.947053	3.190934e-04	-3.599285
49	CAEC=CAEC_Always	0.000000	0.0000000	2.347652	9.877629e-09	-5.732817
50	NCP=1_or_2_meals	15.107914	6.542056	13.886114	5.758978e-12	-6.885488
51	FAVC=FAVC_no	11.340206	3.426791	9.690310	2.185631e-12	-7.022097
52	CH20=More_than_2_liters	15.531915	11.370717	23.476523	4.098730e-20	-9.185434
53	CAEC=CAEC_Frequently	0.000000	0.0000000	9.740260	1.402396e-35	-12.449775
54	CH20=Less_than_1_liter	3.973510	2.803738	22.627373	9.431830e-61	-16.442886
55	FCVC=FCVC_Always	7.052632	10.436137	47.452547	1.968473e-127	-24.014366
56						
57	\$ '3'					
58		Cla/Mod	Mod/Cla	Global	p.value	v.test
59	CAEC=CAEC_Always	100.000000	47.95918	2.347652	7.099873e-68	17.408600
60	CALC=CALC_Frequently	88.135593	53.06122	2.947053	2.045747e-67	17.347904
61	FCVC=FCVC_Never	13.793103	12.24490	4.345654	1.005280e-03	3.289045
62	FAVC=FAVC_no	9.278351	18.36735	9.690310	6.754817e-03	2.708696
63	CH20=1_or_2_liters	6.024096	66.32653	53.896104	1.103054e-02	2.541730
64	CH20=Less_than_1_liter	3.090508	14.28571	22.627373	3.714290e-02	-2.084190
65	FAVC=FAVC_yes	4.424779	81.63265	90.309690	6.754817e-03	-2.708696
66	CALC=CALC_no	2.926829	18.36735	30.719281	4.986010e-03	-2.807936
67	NCP=3_meals	0.000000	0.000000	8.641359	1.128468e-04	-3.861167
68	CALC=CALC_Sometimes	2.108434	28.57143	66.333666	6.227266e-15	-7.799257
69	CAEC=CAEC_Sometimes	2.215909	39.79592	87.912088	5.891423e-32	-11.765303
70						
71	\$ '4'					
72		Cla/Mod	Mod/Cla	Global	p.value	v.test
73	CALC=CALC_no	56.910569	71.283096	30.719281	1.137925e-105	21.832588
74	CAEC=CAEC_Frequently	88.717949	35.234216	9.740260	4.613437e-91	20.237121
75	NCP=1_or_2_meals	69.064748	39.103870	13.886114	4.011970e-66	17.176068

```

76 CH20=Less_than_1_liter 52.759382 48.676171 22.627373 2.122971e-51 15.082133
77 FAVC=FAVC_no 62.371134 24.643585 9.690310 2.140888e-32 11.850430
78 SCC=SCC_yes 66.666667 7.739308 2.847153 1.194836e-11 6.780835
79 FCVC=FCVC_Always 27.263158 52.749491 47.452547 6.917136e-03 2.700807
80 CAEC=CAEC_Always 0.000000 0.000000 2.347652 1.510138e-06 -4.809889
81 FCVC=FCVC_Never 4.597701 0.814664 4.345654 5.037817e-07 -5.024867
82 CALC=CALC_Frequently 0.000000 0.000000 2.947053 4.643107e-08 -5.464463
83 SCC=SCC_no 23.290488 92.260692 97.152847 1.194836e-11 -6.780835
84 CH20=More_than_2_liters 12.978723 12.423625 23.476523 3.202307e-12 -6.968547
85 CH20=1_or_2_liters 17.701576 38.900204 53.896104 1.754237e-14 -7.667469
86 FAVC=FAVC_yes 20.464602 75.356415 90.309690 2.140888e-32 -11.850430
87 NCP=More_than_3_meals 16.892328 53.360489 77.472527 1.546803e-44 -14.000545
88 CAEC=CAEC_Sometimes 18.068182 64.765784 87.912088 3.076376e-62 -16.649012
89 CALC=CALC_Sometimes 10.617470 28.716904 66.333666 2.219414e-88 -19.930364
90
91 > res.HCPC$desc.ind$para
92 Cluster: 1
93 23 111 152 196 228
94 0.1875102 0.1875102 0.1875102 0.1875102 0.1875102
95 -----
96 Cluster: 2
97 106 107 151 169 171
98 0.2381431 0.2381431 0.2381431 0.2381431 0.2381431
99 -----
100 Cluster: 3
101 391 108 202 286 459
102 0.04851977 0.13238048 0.13238048 0.13238048 0.13238048
103 -----
104 Cluster: 4
105 22 53 59 124 251
106 0.2375862 0.2375862 0.2375862 0.2375862 0.2375862
107
108 > #Individuals fareast from the centre
109 > res.HCPC$desc.ind$dist
110 Cluster: 1
111 141 11 51 216 231
112 1.078936 1.078936 1.052467 1.052467 1.052467
113 -----
114 Cluster: 2
115 746 793 894 981 1012
116 1.305863 1.305863 1.305863 1.305863 1.305863
117 -----
118 Cluster: 3
119 67 68 126 415 134
120 3.031465 2.535664 2.535664 2.512287 2.505512
121 -----
122 Cluster: 4
123 199 385 259 328 577
124 1.671027 1.545289 1.519273 1.480808 1.480808

```

2.3 Conclusion

As we have seen, implementing a particular type of Principal Component Analysis helped us understanding the relation between our features, although I was not able to reduce the dimensionality of the data. Various algorithms could have been used, such as Factor Analysis for Mixed Data, or also a classic Principal Component. The former would have been useful in order to detect relations in the data assuming the presence of a latent factor. The latter require to factorize and scale all the data in order to obtain significant results. I ended up preferring Multiple Correspondence Analysis since our data was collected from a survey, a case in which this type of algorithm were purposely implemented.

It is possible to implement the results obtained with the MCA by applying some different type of algorithms,

such as ones that could capture non-linearity or complex relationship between features. Moreover, we are satisfied with the achieved results.

A Code

```
#### Introduction ####
#Libraries and packages required
require(readxl)
require(sandwich)
require(lmtest)
require(expss)
require(MASS)
require(corrplot)
require(dplyr)
require(janitor)
require(data.table)
require(mltools)
require(plyr)
require(caret)
require(tidyr)
require(car)
require(rpart)
require(rpart.plot)
require(tree)
require(partykit)
require(randomForest)
require(nnet)
require(gbm)
require(ggmosaic)
require(grid)
require(gridExtra)
require(parameters)
require(performance)
require(glmnet)
require(Metrics)
require(FactoMineR)
require(factoextra)
require(RColorBrewer)

#Working directory
setwd("~/Desktop/FantaCalcio_Project/Obesity")

#### Data import and cleaning ####
#Data import
data = read.csv(file = "~/Desktop/FantaCalcio_Project/Obesity/data/data.csv")

# 1) Formatting
colnames(data)
data$Age = round(data$Age, 0)
data$Height = round(data$Height, 2)
data$Weight = round(data$Weight, 2)
data$FCVC = round(data$FCVC, 0)
data$NCP = round(data$NCP, 0)
data = data %>%
  mutate(NPC=ifelse(NCP==4,3,NCP))
data$NCP = data$NPC
data = data[, -18]
data$CH20 = round(data$CH20, 0)
data$FAF = round(data$FAF, 0)
data$TUE = round(data$TUE,0)

describe_distribution(data)

#2) Subset DF into categorical variables and numerical variables
```

```

data_num = data[, c("Age", "Height", "Weight")]
data_cha = data[, c("Gender", "family_history_with_overweight", "FAVC", "FCVC",
  "NCP", "CAEC", "SMOKE", "CH20", "SCC", "FAF", "TUE", "CALC", "MTRANS",
  "NObeyesdad")]

#3) Distribution of all numerical variables, represented via Histograms
colnames(data_num)
hist(data_num$Age, main = "Histogram of Age", xlab = "Age", col = "lightblue")
hist(data_num$Height, main = "Histogram of Height", xlab = "Height", col =
  "lightblue")
hist(data_num$Weight, main = "Histogram of Weight", xlab = "Weight", col =
  "lightblue")

#4) Correlation matrix and plot on numerical variables
cor_n = cor(data_num)
corrplot(cor_n, type = "upper", tl.col = "black", tl.srt = 45, title =
  "Correlation plot between numerical variables")

#5) Compare 'NObeyesdad' levels across 'Age', 'Height', 'Weight'
ddply(data, "NObeyesdad", summarize,
  Age = round(mean(Age),2),
  Height = round(mean(Height),2),
  Weight = round(mean(Weight),2))

#6) Distribution of some variables
#ggplot(data, aes(x = Gender, fill = Gender)) +
#  geom_bar(stat = "count") +
#  ggtitle("Distribution for Gender variable") +
#  xlab("Gender") + ylab("Number of records") +
#  theme_bw() +
#  geom_text(aes(label = ..count..), stat = 'count', vjust = -0.4) +
#  labs(fill = "Obesity level") +
#  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1))

#ggplot(data, aes(x = SMOKE, fill = SMOKE)) +
#  geom_bar(stat = "count") +
#  ggtitle("Distribution for Smoking habits response") +
#  xlab("Smokers") + ylab("Number of records") +
#  theme_bw() +
#  geom_text(aes(label = ..count..), stat = 'count', vjust = -0.4) +
#  labs(fill = "Obesity level") +
#  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1))

#ggplot(data, aes(x = FAF, fill = FAF)) +
#  geom_bar(stat = "count") +
#  ggtitle("Distribution for Physical activity frequency response") +
#  xlab("Response") + ylab("Number of records") +
#  theme_bw() +
#  geom_text(aes(label = ..count..), stat = 'count', vjust = -0.4) +
#  labs(fill = "Obesity level") +
#  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1))

#ggplot(data, aes(x = TUE, fill = TUE)) +
#  geom_bar(stat = "count") +
#  ggtitle("Distribution for Time using technology devices response") +
#  xlab("Response") + ylab("Number of records") +
#  theme_bw() +
#  geom_text(aes(label = ..count..), stat = 'count', vjust = -0.4) +
#  labs(fill = "Obesity level") +
#  theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1))

#ggplot(data, aes(x = NObeyesdad, fill = NObeyesdad)) +

```

```

# geom_bar(stat = "count") +
# ggtitle("Distribution across different obesity levels") +
# xlab("Obesity level") + ylab("Number of records") +
# theme_bw() +
# geom_text(aes(label = ..count..), stat = 'count', vjust = -0.4) +
# labs(fill = "Obesity level") +
# theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust=1))

# 6) Distribution of some variables; Gender wise distribution for Weight and Height
#ggplot(data, aes(y = Height, color = Gender)) + #fig3
# geom_boxplot() +
# ggtitle("Gender wise distribution of Height") +
# theme_bw()

#ggplot(data, aes(y = Weight, color = Gender)) + #fig4
# geom_boxplot() +
# ggtitle("Gender wise distribution of Weight") +
# theme_bw()

# 6) Distribution of some variables: Analyzing the preferred Mode of Transport by
  Gender
#ggplot(data) +
# geom_bar(aes(y = MTRANS, fill = Gender),position = position_dodge()) +
# ylab("Mode of Transport") +
# scale_color_manual(values = c("blue", "red"), aesthetics = "fill") +
# ggtitle("Transport preferences by Gender") +
# theme_mosaic()

# 7) Comparing Compare 'NObeyesdad' levels across categorical variables
tabyl(data, NObeyesdad, Gender)
tabyl(data, NObeyesdad, family_history_with_overweight)
tabyl(data, NObeyesdad, FAVC)
tabyl(data, NObeyesdad, FCVC)
tabyl(data, NObeyesdad, NCP)
tabyl(data, NObeyesdad, CAEC)
tabyl(data, NObeyesdad, SMOKE)
tabyl(data, NObeyesdad, CH20)
tabyl(data, NObeyesdad, SCC)
tabyl(data, NObeyesdad, MTRANS)

# 8) Outliers detection
#fig_bp1 = ggplot(data, aes(y = Age, color = Age)) + #fig3
# geom_boxplot() +
# ggtitle("Age") +
# theme_bw()
#fig_bp2 = ggplot(data, aes(y = Height, color = Height)) + #fig3
# geom_boxplot() +
# ggtitle("Height") +
# theme_bw()
#fig_bp3 = ggplot(data, aes(y = Weight, color = Weight)) + #fig3
# geom_boxplot() +
# ggtitle("Weight") +
# theme_bw()
#grid.arrange(fig_bp1, fig_bp2, fig_bp3, ncol = 3)

# 9) Check outliers with Interquartile Range: Percentile: Age column
data_out = data[, c("Age", "Height", "Weight")]
out_summary = as.data.frame(apply(data_out,2,summary))

age_p25 = out_summary[2,1]
age_p75 = out_summary[5,1]

```

```

# 9) Check outliers with Interquartile Range: Percentile: Weight column
wei_p25 = out_summary[2,3]
wei_p75 = out_summary[5,3]

# 9) Check outliers with Interquartile Range: Interquartile Ranges
IR_age = age_p75 - age_p25
IR_wei = wei_p75 - wei_p25

# 9) Check outliers with Interquartile Range: Determine the lower and upper limit
lower_lim_age = age_p25 - (1.5 * IR_age)
upper_lim_age = age_p75 + (1.5 * IR_age)

lower_lim_weight = wei_p25 - (1.5 * IR_wei)
upper_lim_weight = wei_p75 + (1.5 * IR_wei)

# 9) Check outliers with Interquartile Range: Dropping observations which do not
      satisfy IQR outliers approach
data_iqr = data
data_iqr = data_iqr[!(data_iqr$Age > upper_lim_age | data_iqr$Age <
      lower_lim_age),]
data_iqr = data_iqr[!(data_iqr$Weight > upper_lim_weight | data_iqr$Weight <
      lower_lim_weight),] #same outliers for Weight and Height

# 9) Check outliers with Interquartile Range: boxplot without outliers
#fig_bp4 = ggplot(data_iqr, aes(y = Age, color = Age)) + #fig3
# geom_boxplot() +
# ggtitle("Age") +
# theme_bw()
#fig_bp5 = ggplot(data_iqr, aes(y = Height, color = Height)) + #fig3
# geom_boxplot() +
# ggtitle("Height") +
# theme_bw()
#fig_bp6 = ggplot(data_iqr, aes(y = Weight, color = Weight)) + #fig3
# geom_boxplot() +
# ggtitle("Weight") +
# theme_bw()
#grid.arrange(fig_bp4, fig_bp5, fig_bp6, ncol = 3)

#rm(data_cha, data_num, age_p25, age_p50, age_p75, hei_p25, hei_p50, hei_p75,
      wei_p25, wei_p50, wei_p75, IR_age, IR_wei, IR_wei,
#      lower_lim_age, lower_lim_height, lower_lim_weight, upper_lim_age,
      upper_lim_height, upper_lim_weight, outliers,
#      data_out, cor_n, out_summary, find_outliers_age, find_outliers_weight)
#rm(fig_bp1, fig_bp2, fig_bp3, fig_bp4, fig_bp5, fig_bp6)

#### Data preparation for regression\classification ####
# 1) Data Pre-Processing
sapply(data_iqr, class)

# 2) one-hot-encoding for binary categorical variables
data_cha_ohe = data_iqr[, c("Gender", "family_history_with_overweight", "FAVC",
      "SMOKE", "SCC", "MTRANS")]
dmy = dummyVars("~", data = data_cha_ohe)
dat_transformed = data.frame(predict(dmy, newdata = data_cha_ohe))

data_iqr = cbind(data_iqr, dat_transformed)
data_iqr = data_iqr[, -c(1, 5, 6, 10, 12, 16)]
rm(data_cha_ohe, dmy, dat_transformed)

# 3) Ordinal Encoding: used for ordinal categorical variables "CAEC", "CALC",
      "NObyesdad"

```

```

data_iqr$CAEC = factor(data_iqr$CAEC, levels = c("no", "Sometimes", "Frequently",
  "Always"), labels = c(1, 2, 3, 4))
data_iqr$CALC = factor(data_iqr$CALC, levels = c("no", "Sometimes", "Frequently",
  "Always"), labels = c(1, 2, 3, 4))
data_iqr$NObeyesdad = factor(data_iqr$NObeyesdad, levels =
  c("Insufficient_Weight", "Normal_Weight", "Overweight_Level_I",
  "Overweight_Level_II", "Obesity_Type_I", "Obesity_Type_II",
  "Obesity_Type_III"), labels = c(0:6))

col_order = c("GenderFemale", "GenderMale", "Age", "Height", "Weight",
  "family_history_with_overweightno", "family_history_with_overweightyes",
  "FAVCno", "FAVCyes", "FCVC", "NCP", "CAEC",
  "SMOKEno", "SMOKEyes", "CH20", "SCCno", "SCCyes", "FAF", "TUE",
  "CALC", "MTRANSAutomobile", "MTRANSBike", "MTRANSMotorbike",
  "MTRANSPublic_Transportation",
  "MTRANSWalking", "NObeyesdad")
data_iqr = data_iqr[, col_order]
describe_distribution(data_iqr)

dev.off()
boxplot(data_iqr)
title("Comparing_boxplot()s")

# 4) Scaling variables
data_iqr$CAEC = as.numeric(levels(data_iqr$CAEC))[data_iqr$CAEC]
data_iqr$CALC = as.numeric(levels(data_iqr$CALC))[data_iqr$CALC]
data_iqr$NObeyesdad = as.numeric(levels(data_iqr$NObeyesdad))[data_iqr$NObeyesdad]
data_iqr_2 = scale(x = as.data.table(data_iqr), center = TRUE, scale = TRUE)

boxplot(data_iqr_2)
title("Comparing_boxplot()s_scaled")
summary(data_iqr_2)
colnames(data_iqr)
data_iqr = data_iqr[,-c(2, 6, 8, 13, 16, 25)]

### 5) DF with only categorical variable scaled and numerical ones unchanged
scaled = as.data.frame(scale(x = as.data.table(data_iqr[,-c(2:4, 20)]), center =
  TRUE, scale = TRUE))

true = data_iqr[, c(2:4, 20)]
sc_data = cbind(scaled, true)

col_order = c("GenderFemale", "Age", "Height", "Weight",
  "family_history_with_overweightyes",
  "FAVCyes", "FCVC", "NCP", "CAEC",
  "SMOKEyes", "CH20", "SCCyes", "FAF", "TUE",
  "CALC", "MTRANSAutomobile", "MTRANSBike", "MTRANSMotorbike",
  "MTRANSPublic_Transportation",
  "NObeyesdad")
sc_data = sc_data[, col_order]
rm(true, scaled)

#### before OLS REGRESSION ####
# Training and Test splitting set
set.seed(123)
split_train_test = createDataPartition(y = sc_data$NObeyesdad, p=0.8, list = F)
train_data = sc_data[split_train_test ,]
test_data = sc_data[-split_train_test ,]
dim(test_data)
dim(train_data)
#### 0) Part 1: LINEAR REGRESSION ####
lin_mod = lm(formula = NObeyesdad ~ ., data = train_data)

```



```

lin_sum = summary(lin_mod)

par(mfrow = c(2,2))
plot(lin_mod)
dev.off()

check_model(lin_mod)
check_normality(lin_mod)
check_heteroscedasticity(lin_mod)
check_autocorrelation(lin_mod)
check_collinearity(lin_mod)

coeftest(lin_mod, vcov. = vcovHC, type = "HC1")
lin_mod_performance = performance(lin_mod)

# STEPWISE REGRESSION
lin_mod_step = lm(formula = NObeyesdad ~ ., data = train_data)
lin_mod_step = select_parameters(lin_mod_step)
lin_step_sum = summary(lin_mod_step)

coeftest(lin_mod_step ,vcov. = vcovHC , type = "HC1")

check_model(lin_mod_step)
check_normality(lin_mod_step)
check_heteroscedasticity(lin_mod_step)
check_autocorrelation(lin_mod_step)
check_collinearity(lin_mod_step)

lin_step_performance = performance(lin_mod_step)

# LASSO REGRESSION
x = as.matrix(train_data[, -20])
y = train_data$NObeyesdad
cv_lasso = cv.glmnet(x,y)
plot(cv_lasso)

coef = coef(cv_lasso , s = cv_lasso$lambda.min)
coefname = coef@Dimnames [[1]][ -1]
coef = coefname[coef@i]
coef

fmla = as.formula(paste("y~", paste(coef, collapse = "+")))

lin_lasso = lm(fmla, data=train_data)
lin_lasso_sum = summary(lin_lasso)
check_model(lin_lasso)

check_normality(lin_lasso)
check_heteroscedasticity(lin_lasso)
check_autocorrelation(lin_lasso)
check_collinearity(lin_lasso)

coeftest(lin_lasso, vcov. = vcovHC, type='HC1')

lin_lasso_performance = performance(lin_lasso)

# compare linear regressions
compare_performance(lin_mod, lin_lasso, lin_mod_step, rank = T)
plot(compare_performance(lin_mod, lin_lasso, lin_mod_step, rank = T))

# Robust lasso estimator
lin_mod_lasso_robust = rlm(fmla, data=train_data, psi = psi.bisquare)

```

```

hweights = data.frame(resid = lin_mod_lasso_robust$resid, weight =
  lin_mod_lasso_robust$w)
hweights2 = hweights[order(lin_mod_lasso_robust$w),]
hweights2 [1:10 ,]

rob_se_pan = list(sqrt(diag(vcovHC(lin_mod , type = "HC1"))),
  sqrt(diag(vcovHC(lin_mod_step , type = "HC1"))),
  sqrt(diag(vcovHC(lin_lasso , type = "HC1"))),
  sqrt(diag(vcovHC(lin_mod_lasso_robust , type = "HC1"))))

#### 1) Part 1: MULTINOMIAL LOGISTIC REGRESSION ####
train_data_log = train_data
test_data_log = test_data
train_data_log$NObeyesdad = factor(train_data_log$NObeyesdad, levels = c(1:7),
  labels = c("Insufficient_Weight", "Normal_Weight", "Overweight_Level_I",
  "Overweight_Level_II", "Obesity_Type_I", "Obesity_Type_II",
  "Obesity_Type_III"))
test_data_log$NObeyesdad = factor(test_data_log$NObeyesdad, levels = c(1:7),
  labels = c("Insufficient_Weight", "Normal_Weight", "Overweight_Level_I",
  "Overweight_Level_II", "Obesity_Type_I", "Obesity_Type_II",
  "Obesity_Type_III"))

# Setting the 'reference' level
train_data_log$NObeyesdad = relevel(train_data_log$NObeyesdad, ref =
  "Insufficient_Weight")
multinom_model = multinom(NObeyesdad ~ ., data = train_data_log) #training the
  model
summary(multinom_model)
exp(coef(multinom_model)) # Convert the coefficients to odds by taking the
  exponential of the coefficients.
head(round(fitted(multinom_model), 2))# The predicted values are saved as
  fitted.values in the model object.

# Predicting & Validating the model: Predicting the values for train dataset
train_data_log$ClassPredicted = predict(multinom_model, newdata = train_data_log,
  "class")
tab = table(train_data_log$NObeyesdad, train_data_log$ClassPredicted) # Building
  classification table
tab
round((sum(diag(tab))/sum(tab))*100,2)

# Predicting the class on test dataset.
test_data_log$ClassPredicted = predict(multinom_model, newdata = test_data_log,
  "class")
tab_1 = table(test_data_log$NObeyesdad, test_data_log$ClassPredicted)# Building
  classification table
tab_1
round((sum(diag(tab_1))/sum(tab_1))*100,2)

#### 2) Part 1: REGRESSION DECISION TREES####
train_data_log = train_data
test_data_log = test_data
train_data_log$NObeyesdad = factor(train_data_log$NObeyesdad, levels = c(1:7),
  labels = c("Insufficient_Weight", "Normal_Weight", "Overweight_Level_I",
  "Overweight_Level_II", "Obesity_Type_I", "Obesity_Type_II",
  "Obesity_Type_III"))
test_data_log$NObeyesdad = factor(test_data_log$NObeyesdad, levels = c(1:7),
  labels = c("Insufficient_Weight", "Normal_Weight", "Overweight_Level_I",
  "Overweight_Level_II", "Obesity_Type_I", "Obesity_Type_II",
  "Obesity_Type_III"))

```

```

tree_1 = rpart(NObeyesdad~., data = train_data_log, method = "class")
printcp(tree_1)
summary(tree_1)

cv_tree_1 = plotcp(tree_1) #size of the tree = 10
rpart.plot(tree_1)

prune_tree_1 = prune(tree_1,cp=0.023)
rpart.plot(prune_tree_1)

#prediction on train data
dt_1_yhat = as.numeric(predict(tree_1, newdata=train_data_log, type= "class"))
dt_1_train = as.numeric(train_data_log[, "NObeyesdad"]) #output vector in test data
plot(dt_1_yhat,dt_1_train);abline(0,1, col = "red") #how is the distribution
    between the final units in the nodes

# Prediction on test data
dt_1_pred = as.numeric(predict(tree_1, test_data, type= "class"))

# Confusion Matrix
confMat = table(test_data[, "NObeyesdad"],dt_1_pred)
dt_1_accuracy = sum(diag(confMat))/sum(confMat)

TREE = c(rmse(test_data$NObeyesdad, dt_1_pred),
          mae(test_data$NObeyesdad, dt_1_pred),
          R2(train_data$NObeyesdad, dt_1_yhat))

tab_2 = table(test_data_log$NObeyesdad,dt_1_pred)
tab_accuracy = sum(diag(tab_2))/sum(tab_2)

#### 3) Part 1: RANDOM FOREST REGRESSION ####
rf_1 = randomForest(NObeyesdad ~ .,data=train_data, mtry = sqrt(19), type =
    "regression")
plot(rf_1, main = "Random_Forest_Error")
rf_1_y_hat = as.numeric(predict(rf_1 ,newdata=test_data))
test_data$NObeyesdad = as.numeric(test_data$NObeyesdad)
mean((rf_1_y_hat - test_data$NObeyesdad)^2)

ggplot() +
  geom_point(aes(x = test_data$NObeyesdad, y = rf_1_y_hat)) +
  geom_abline()

impo = round(importance(rf_1),2)
varImpPlot (rf_1)

oob.err=double(19)
test.err=double(19)
for(mtry in 1:19){
  fit=randomForest(NObeyesdad~.,data=train_data,mtry=mtry,ntree=400)
  oob.err[mtry]=fit$mse[400]
  pred=predict(fit,test_data)
  test.err[mtry]=with(test_data,mean((NObeyesdad-pred)^2))
  cat(mtry," ")
}
oob.err
test.err
matplot(1:mtry,cbind(test.err,oob.err),pch=19,col=c("red","blue"),type="b",ylab="Mean_
    Squared_Error");
  legend("topright",legend=c("Test","OOB"),pch=19,col=c("red","blue"))

RF = c(rmse(test_data$NObeyesdad,rf_1_y_hat),
        mae(test_data$NObeyesdad, rf_1_y_hat),

```

```

R2(test_data$NObeyesdad, rf_1_y_hat))

#### Part 1: COMPARISONS ####
test_pred = predict(lin_mod, newdata = test_data)
train_pred = predict(lin_mod, newdata = train_data)
OLS = c(rmse(test_data$NObeyesdad, test_pred),
        mae(test_data$NObeyesdad, test_pred),
        R2(train_data$NObeyesdad, train_pred))
test_pred = predict(lin_mod_step, newdata = test_data)
train_pred = predict(lin_mod_step, newdata = train_data)
Stepwise = c(rmse(test_data$NObeyesdad, test_pred),
             mae(test_data$NObeyesdad, test_pred),
             R2(train_data$NObeyesdad, train_pred))
test_pred = predict(lin_lasso, newdata = test_data)
train_pred = predict(lin_lasso, newdata = train_data)
LASSO = c(rmse(test_data$NObeyesdad, test_pred),
          mae(test_data$NObeyesdad, test_pred),
          R2(train_data$NObeyesdad, train_pred))
test_pred = predict(lin_mod_lasso_robust, newdata = test_data)
train_pred = predict(lin_mod_lasso_robust, newdata = train_data)
Robust = c(rmse(test_data$NObeyesdad, test_pred),
           mae(test_data$NObeyesdad, test_pred),
           R2(train_data$NObeyesdad, train_pred))
metrics = data.frame(OLS, Stepwise, LASSO, Robust, TREE, RF)
row.names(metrics) = c("RMSE", "MAE", "R2")
metrics
#### 1) Part 2: INTRODUCTION ####
rm(list = ls())
dev.off()
data = read.csv(file = "~/Desktop/FantaCalcio_Project/Obesity/data/data.csv")

# Data cleaning
data$Age = round(data$Age, 0)
data$Height = round(data$Height, 2)
data$Weight = round(data$Weight, 2)
data$FCVC = round(data$FCVC, 0)
data$NCP = round(data$NCP, 0)
data = data %>%
  mutate(NPC=ifelse(NCP==4,3,NCP))
data$NCP = data$NPC
data = data[, -18]
data$CH20 = round(data$CH20, 0)
data$FAF = round(data$FAF, 0)
data$TUE = round(data$TUE, 0)

# Replace nominal variables with categorical variables
data["FCVC"][data["FCVC"] == 1] = "Never"
data["FCVC"][data["FCVC"] == 2] = "Sometimes"
data["FCVC"][data["FCVC"] == 3] = "Always"

data["NCP"][data["NCP"] == 1] = "1_or_2_meals"
data["NCP"][data["NCP"] == 2] = "3_meals"
data["NCP"][data["NCP"] == 3] = "More_than_3_meals"

data["CH20"][data["CH20"] == 1] = "Less_than_1_liter"
data["CH20"][data["CH20"] == 2] = "1_or_2_liters"
data["CH20"][data["CH20"] == 3] = "More_than_2_liters"

data["FAF"][data["FAF"] == 0] = "No_activity"
data["FAF"][data["FAF"] == 1] = "1_or_2_days"
data["FAF"][data["FAF"] == 2] = "2_or_3_days"
data["FAF"][data["FAF"] == 3] = "4_or_5_days"

```

```

data["TUE"][data["TUE"] == 0] = "0-to_2_hours"
data["TUE"][data["TUE"] == 1] = "3_to_5_hours"
data["TUE"][data["TUE"] == 2] = "More_than_5_hours"

quanti = data[,c(2,3,4)]
quali = data[, c(1, 5:16)]
quali = as.data.frame(unclass(quali), stringsAsFactors = TRUE)
data = cbind(quanti, quali)
rm(quali, quanti)
#### 2) Part 2: MCA ####
#1) Model
data_active = data[, c(6:9, 11, 12, 15)]
par(mfrow = c(2, 4))
for (i in 1:7) {
  plot(data_active[,i], main=colnames(data_active)[i],
        ylab = "Count", col="beige", las = 2)
}

res.mca = MCA(data_active, graph = TRUE, ncp = 3)

#Outliers detection with Factoinvestigate and removele
outliers = FactoInvestigate::outliers(res.mca)
outliers$ID

data_active = data_active[-outliers$ID, ]

res.mca = MCA(data_active, graph = TRUE, ncp = 3)

#2) Visualization and Interpretation: Eigenvalues - Variance
eig.val <- get_eigenvalue(res.mca)
fviz_screplot(res.mca, addlabels = TRUE, ylim = c(0, 45))

#3) Bi-plot
fviz_mca_biplot(res.mca,
                 repel = TRUE, # Avoid text overlapping (slow if many point)
                 ggtheme = theme_minimal())

#4) Correlation between variables and principal dimensions
fviz_mca_var(res.mca, choice = "mca.cor",
              repel = TRUE, # Avoid text overlapping (slow)
              ggtheme = theme_minimal())

#5) extract the results for variable categories
var = get_mca_var(res.mca)

#6) coordinates of variables to create a scatter plot
round(var$coord, 2)

#7) contains the contributions (in percentage) of the variables to the definition
    of the dimensions.
round(var$contrib, 2)

#8) represents the quality of the representation for variables on the factor map.
round(var$cos2, 2)

#9) Color by cos2 values: quality on the factor map:
fviz_mca_var(res.mca, col.var = "cos2",
              gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
              repel = TRUE, # Avoid text overlapping
              ggtheme = theme_minimal())

```

```

#10) Graph of individuals
ind = get_mca_ind(res.mca)

#11) Coordinates of column points
head(ind$coord)

#12) Quality of representation
head(ind$cos2)

#13) Contributions
head(ind$contrib)

#14) Color individuals by groups
fviz_mca_ind(res.mca,
             label = "none", # hide individual labels
             habillage = "FAVC", # color by groups
             palette = display.brewer.pal(n = 7, name = 'RdBu'),
             addEllipses = TRUE, ellipse.type = "confidence",
             ggtheme = theme_minimal())

#15) Dimension description
res.desc = dimdesc(res.mca, axes = c(1,2))

#16) Description of dimension 1
res.desc[[1]]

#17) Description of dimension 2
res.desc[[2]]

#### 4) Part 2: HCPC ####
res.HCPC=HCPC(res.mca, graph = FALSE)

#1) Dendrogram
fviz_dend(res.HCPC, show_labels = FALSE)

#2) Individuals factor map
fviz_cluster(res.HCPC, geom = "point", main = "Factor_map")

#3) Description by variables
head(res.HCPC$desc.var$test.chi2,10)

#4) Description by variable categories
head(res.HCPC$desc.var$category,10)

#5) Description by principal components
res.HCPC$desc.axes

#6) Individuals closest to the centre
res.HCPC$desc.ind$para

#7) Individuals farthest from the centre
res.HCPC$desc.ind$dist

#8) Plots
plot.HCPC(res.HCPC,choice='tree',title='Hierarchical_tree')
plot.HCPC(res.HCPC,choice='map',draw.tree=FALSE,title='Factor_map')
plot.HCPC(res.HCPC,choice='3D.map',ind.names=FALSE,centers.plot=FALSE,angle=60,title='Hierarchical_map')

```